

IndiaLens

Quantitative Education & Career Intelligence Operating System
Product Architecture & System Assembly Blueprint

DOCUMENT STANDARD	PRODUCTION ARCHITECTURE	INTELLIGENCE STACK	DOCUMENT STATUS
Industry-Grade Engineering Spec	Next.js 14 / FastAPI / PG15	DeepSeek V4 Flash + Sonar + Quant ML	Approved Engineering Standard

TABLE OF CONTENTS · MASTER SPECIFICATIONS INDEX

01 Executive Product Brief & Foundational Thesis	02 Global & Domestic Competitive Landscape
03 Curated Course Marketplace & Dynamic Affiliate Engine	04 The Student Asset Pricing Engine: Mathematical Framework
05 The Flagship AI Risk & Future Labor Demand Index	06 The 3PL IRT Adaptive Psychometric Engine (CAT)
07 Machine Learning Model Specifications & Pipelines	08 The Student Intelligence Layer: Dual-Engine Architecture (DeepSeek + Sonar)
09 Caliber-Calibrated Admissions Portfolio Engine	10 Global Degree ROI & Cross-Border Actuarial Models
11 Global Portfolio Builder & Admissions Spike System	12 Data Ingestion Infrastructure & 18 API Connectors
13 Production Relational Database Schema (21 Tables)	14 Page-by-Page UI/UX Architecture & Screen Specs
15 Commercialization, Unit Economics & 3-Year Projections	

SECTION 01 · STRATEGIC FOUNDATION

Executive Product Brief & Foundational Thesis

IndiaLens is built as India's first sovereign, quantitative Education & Career Intelligence Operating System. For decades, college and career decisions in India and emerging markets have been governed by name-brand hype, unverified rumors, intense family anxiety, and opaque marketing billboards. Students and parents routinely commit substantial portions of their family savings—often taking high-interest, non-forgivable education loans—to pay for college degrees without knowing their true expected financial return, the real risk of struggling to repay their loan, or whether AI tools will automate that job before they even graduate.

💡 IN PLAIN ENGLISH · WHY INDIALENS EXISTS

Choosing a college is the single largest financial and emotional bet an Indian family makes. Today, parents spend ₹10 Lakhs to ₹1.2 Crores based on glossy brochures, coaching institute advertisements, and relatives' advice. **Traditional rankings rank the college; IndiaLens evaluates the student inside the college.** A degree is not a decorative luxury or an automatic guarantee of wealth—it is an investment whose true financial payoff depends on your individual abilities, your family budget, and whether the job market will still need those skills ten years from now. IndiaLens replaces panic and marketing hype with honest, mathematical reality.

THE FOUNDATIONAL AXIOM

Rankings measure institutions; **IndiaLens measures the student within the institution.** An educational degree is not a luxury consumer good or an immutable social stamp; it is a long-duration, highly leveraged, stochastic capital asset whose financial and developmental return is strictly conditioned on the individual student's latent traits, financial constraints, and macroeconomic tailwinds.

1.1 Target Audience Personas & Critical User Needs

The platform serves four distinct user groups, giving each of them the exact clarity and protection they need:

- **The Aspirational Student (Ages 16–22):** Preparing for high-stakes competitive entrance exams (JEE, NEET, CUET, CAT, GRE, GMAT, SAT) under heavy mental pressure. They need honest, personalized guidance that cuts through the toxic "IIT-or-failure" mentality. IndiaLens builds them a smart 4-tier college portfolio (Dream, Reach, Target, Safety) tuned specifically to their actual strengths and testing stamina.
- **The Anxious Parent / Family Sponsor:** The primary financial investor facing massive out-of-pocket costs (from ₹10 Lakhs for domestic engineering up to ₹1.2 Crores for international degrees). They need clear, easy-to-understand balance sheet stress tests, plain loan EMI projections under bad economies, and realistic take-home monthly salary estimates so their family's financial security is never put in jeopardy.
- **The High School Counselor & Coaching Educator:** Mentors looking for a professional, data-backed copilot to replace generic advice spreadsheets. They get multi-year student tracking, objective strength-testing tools, and verified salary databases.
- **The Education Lender & Enterprise Recruiter (B2B):** Non-Banking Financial Companies (NBFCs), banks, loan underwriters, and hiring managers who need reliable, algorithmic risk scores to verify student employability and loan repayment safety with precision.

1.2 Systemic Market Failures in the Status Quo


Today's educational counseling market suffers from four severe, hidden conflicts of interest that trap innocent families:

TABLE 1.1 – ANALYSIS OF INCUMBENT SYSTEMIC FAILURES

Current Method	What Is Broken Behind the Scenes	How It Harms Students & Families
Official Rankings (NIRF, QS, Times)	Rankings prioritize vanity metrics: campus land area, spending on buildings, patents filed (even if useless), and faculty PhD ratios—completely ignoring real median graduate salaries and job placement consistency.	Students pay massive fees for fancy campuses while average graduate salaries fail to even beat inflation.
College Portals (Shiksha, CollegeDunia)	The "Lead-Generation" Broker Model: Private colleges pay these portals ₹300 to ₹2,500 every time they sell a student's phone number and email. The search results secretly push whichever low-quality private college pays the highest advertising bounty.	Students are relentlessly pushed toward expensive, sub-par private colleges that hide poor placement records behind sponsored search banners.
Study-Abroad Agents & Consultants	The Commission Broker Model: Overseas agents get a secret 15% to 30% cash kickback from foreign colleges on the student's first-year tuition fee.	Students are steered into low-tier diploma mills in regional Canada, the UK, or Australia with almost zero chance of landing a work visa or staying after graduation.
Blindness to AI Job Automation	Counselors quote 5-year-old salary records as if they are guaranteed forever, ignoring how modern AI tools are already wiping out junior coding, routine accounting, and basic data entry roles.	Graduates spend 4 years earning a degree only to discover that the entry-level tasks they trained for have already been automated by AI.

1.4 The Operating System Blueprint: 5-Stage System Assembly Flow

To keep the platform modular, transparent, and easy to understand, IndiaLens operates like a clean 5-stage precision assembly line where data moves smoothly from verified facts to actionable student advice:



Global & Domestic Competitive Landscape

The education guidance market is divided between commission-driven college brokers in India and static government tables abroad. To give students an uncompromised advantage, IndiaLens combines mathematical depth with 100% student-aligned honesty, outperforming every existing tool across all major categories.

TABLE 2.1 – COMPREHENSIVE COMPETITIVE BENCHMARK (TOP GLOBAL & DOMESTIC PLAYERS)

Evaluation Dimension	IndiaLens	Shiksha.com (India)	CollegeDunia (India)	Careers360 (India)	Niche.com (US)	College Scorecard (US)
Business Model	Independent Student OS & Curated Upskilling	College Lead Gen & Paid Banners	Cost-Per-Lead (CPL) Brokerage	Media Ads & Basic Cutoff Predictors	Digital Ads & Lead Generation	Public Government Table
What They Actually Evaluate	The Individual Student (Traits, budget, talent)	The College Brand	College Profile & Reviews	Single Exam Rank Cutoffs	Campus Review Ratings	College Aggregate Averages
How Success is Measured	20-Year Net Cash Return, Breakeven & Loan Default Risk	Self-Reported Unverified Placement	User Star Ratings	Exam Rank Match Probability	Campus Vibe Grades (A+ to F)	10-Year Median Debt & Earnings
Personalization	Smart Adaptive Quizzing (3PL IRT) + Career Roadmap Graph	None (Only Stream & City filters)	Basic Fee Sliders	Basic Cutoff Match	Profile Search Filters	None (One size fits all)
AI Threat Analysis	Live 8-Vector AI Risk Index & Task Automation Curves	None (Assumes jobs never change)	None (Zero risk modeling)	None (No automation data)	None (No job task links)	None (Outdated tax records)
Loan Stress Testing	10,000-Future Monte Carlo Simulation with Rate Changes	None (Only shows tuition)	Simple Fixed-Rate Calculator	Basic Fee Table	Simple Net Price Calculator	Static Median Debt at Exit
Global Degrees & Visas	Cross-Border Net Return with Visa Lottery Odds & FX Math	Lead capture for study abroad agents	Overseas coaching ads	Static articles	US Domestic Only	US Domestic Only
Data Freshness	Real-Time & Weekly Automated Sync (Airflow)	Annual brochure uploads	Irregular user reviews	Once a year post-exam	Annual survey cycle	2 to 3-year lag
Advice Loyalty	100% Student Fiduciary (No College Ads Allowed)	College marketing catalog	Crowdsourced comments	Editorial magazine style	Sponsored social reviews	Neutral government data

2.1 IndiaLens's Structural Moat (Why Competitors Cannot Copy Us)

IndiaLens's competitive edge rests on three powerful structural advantages that traditional college directories cannot replicate without destroying their own business:

- **The Conflict-of-Interest Moat:** Incumbent portals get over 85% of their money directly from college marketing departments. They cannot afford to show a college's real negative ROI or a high AI job risk without getting sued or losing their biggest advertisers. IndiaLens is funded by students and lenders, making radical honesty our greatest sales advantage.
- **The Mathematical & Actuarial Engine:** Competitors only show static numbers. IndiaLens runs dynamic future salary predictions, adaptive psychometric talent tests, and 10,000-scenario loan stress simulations that require deep quantitative engineering.
- **The Lifelong Student Ledger:** Unlike search portals where users visit once and leave, IndiaLens acts as a persistent personal career advisor—remembering your strengths, budget, and coursework from 11th grade through your first job promotion.

Curated Course Marketplace & Dynamic Affiliate Engine

Instead of displaying random, annoying pop-up ads, IndiaLens features a curated upskilling recommender. When our algorithms notice that a student's chosen college degree has a weak syllabus, an emerging AI automation threat, or a missing skill needed for high-paying roles, it recommends top-tier, verified courses and certifications to bridge that specific gap.

💡 IN PLAIN ENGLISH • HOW WE RECOMMEND UPSKILLING COURSES

We never push courses just because an advertiser pays us. A course is only suggested to you if it passes four strict checks: (1) **Does it teach skills you actually need and don't already have?** (2) **Does it directly help you reach your 10-year target career?** (3) **Does it teach high-level problem solving that AI cannot easily automate?** and (4) **Can your family comfortably afford it?** If a course fails this test or scores below 75%, it is completely blocked from your screen.

3.1 Mathematical Course Match & Relevance Formulation

Every candidate course c evaluated for student s must surpass a rigorous Match Threshold score of $\mathcal{M}(s, c) \geq 0.75$ before entering the recommendation pool:

EQUATION 3.1 – DYNAMIC COURSE MATCH SCORING TENSOR

$$\mathcal{M}(s, c) = w_1 \cdot \Delta_{\text{skill}}(s, c) + w_2 \cdot \text{Align}_{\text{career}}(s, c) + w_3 \cdot \rho_{\text{AI}}(c) + w_4 \cdot \mathcal{B}_{\text{fit}}(s, c)$$

Breaking Down the Four Core Factors in Plain Terms:

- $\Delta_{\text{skill}}(s, c) = 1 - \frac{|\mathbf{k}_s \cap \mathbf{k}_c|}{|\mathbf{k}_c|}$: **The Missing Skill Factor.** Checks what percentage of the course curriculum teaches skills you currently lack. If a course spends 80% of its time on things you already know, its score is heavily penalized.
- $\text{Align}_{\text{career}}(s, c) = \cos(\mathbf{v}_{\text{target}}(s), \mathbf{u}_{\text{syllabus}}(c))$: **Career Goal Alignment.** Measures how closely the skills taught in the course match the real-world requirements of your dream 10-year job.
- $\rho_{\text{AI}}(c) = \sum_{j \in \mathbf{k}_c} \text{ResilienceScore}(j) \cdot \text{Weight}(j)$: **The AI Defense Factor.** Rewards courses that teach deep systems thinking, complex problem-solving, and hands-on execution, while downranking courses that only teach routine syntax or rote formulas that AI can easily do.
- $\mathcal{B}_{\text{fit}}(s, c) = \max\left(0, 1 - \frac{\text{Price}_{\text{enr}}(c)}{\text{DiscretionaryMonthlyBudget}(s)}\right)$: **Budget Accessibility.** Prevents the system from recommending overpriced courses to students with modest family budgets, steering them instead toward high-yield certifications, financial aid, or free open-source equivalents.
- Factor Weights:** Initialized to: $w_1 = 0.35$ (Skills needed), $w_2 = 0.25$ (Career fit), $w_3 = 0.25$ (AI safety), $w_4 = 0.15$ (Budget fit), summing to 1.0.

ENGINEERING INTEGRATION GUIDE • COURSE MATCH EXECUTION FLOW

Input Pipeline: Triggered during client dashboard load (`GET /api/v2/marketplace/recommended`). Ingests the student's latent profile from `student_reports.profile_data` and active target node from `personal_intelligence.path_graph`. Fetches candidate items from `course_marketplace` table.

Runtime & Execution Budget: Vectorized matrix multiplication in Python backend executing in < 25ms. Drops all candidates where $\mathcal{M}(s, c) < 0.75$. Top 6 ranked courses returned with signed affiliate redirect links.

Telemetry & Destination: Impressions and conversions logged asynchronously to `course_impressions` table for commercial attribution and conversion tracking.

3.2 Categorical Sourcing & Affiliate API Integration

The marketplace connects directly to global education providers across five verified categories:

TABLE 3.1 – COURSE MARKETPLACE SOURCING & AFFILIATE INTEGRATION MATRIX

Category Vertical	Strategic Providers	Integration Architecture	Attribution Window	Affiliate Economics
Technical Upskilling & Micro-Degrees	Coursera, edX, Udacity, DeepLearning.AI	REST API Server-to-Server Webhook + Deeplink Tokenization	30-Day Cookie / Universal User UUID	15% – 45% Net Revenue Share per subscription / degree enrollment
Vocational Coding & Applied Engineering	Udemy, DataCamp, Pluralsight, Educative	Impact.com & Rakuten LinkShare Gateway	30-Day Cookie Window	20% – 50% Commission on standalone transactional purchases
Competitive Exam Preparation	Unacademy, Physics Wallah, Allen Digital	Direct OAuth Enterprise Partner Webhooks	Direct Student Lead Attribution	₹1,500 – ₹8,000 Flat Bounty per paid batch enrollment
Global Language & Relocation Mastery	Duolingo English Test, IELTS Official, Goethe-Institut	Direct Developer Affiliate API Links	45-Day Conversion Window	\$15 – \$40 CPA per verified exam / preparation voucher purchase
Cloud & Enterprise Certifications	AWS Training, Google Cloud Skills Boost, CFA Institute	Direct Voucher Sourcing API	Instant Redemption	8% – 15% Margin on certification voucher bundles

3.3 Seller Auction & Contextual Placement Policy

Course creators and online educators may participate in our recommendation auction under three strict rules that protect students:

- Quality Gatekeeper:** If a course does not achieve an organic relevance match score of $\mathcal{M}(s, c) \geq 0.75$ for that specific student, it cannot be shown at all. Money cannot buy visibility for a bad or irrelevant course.
- Transparent Placement Score:** For courses that pass the 75% quality bar, recommendations are ranked by: $\text{RankScore} = \mathcal{M}(s, c) \times \text{Bid}_{\text{CPC}}$. High-relevance courses beat high-spending low-relevance courses every time.
- Honest Salary Uplift Display:** Every course card must explicitly state its estimated real salary impact: $\Delta \hat{S}_{Y1} = \text{ModelUplift}(\mathbf{k}_c)$, calculated by comparing the real career earnings of graduates who took that certificate versus those who did not.

SECTION 04 • ECONOMETRIC MODELING

The Student Asset Pricing Engine: Complete Mathematical Framework

The core philosophy of **IndiaLens** is that every college degree or career pathway is a multi-decade financial investment. Instead of relying on misleading college brochures that quote cherry-picked starting salaries, IndiaLens models your real future earnings using probability-adjusted income taxes, loan repayment schedules, economic recessions, and multi-year simulations.

💡 IN PLAIN ENGLISH • THE BIG QUESTION WE ANSWER

"If my family invests ₹20 Lakhs in this college today, how much real money will I actually take home over the next 20 years?"

Most college counselors only look at the highest salary on a brochure. IndiaLens calculates your real take-home cash. We subtract taxes, deduct monthly loan EMI, adjust for inflation (rising living costs), simulate

the chance of recessions or layoffs, calculate whether AI will replace that job, and if studying abroad, calculate visa lottery odds and currency exchange rates. This gives you one clear rupee amount: *is this college an investment that makes you financially secure, or an overpriced trap?*

4.1 The Master Net Present Value (NPV) Formulation

The 20-Year Net Present Value represents the total real-world wealth generated by your college education compared to working right away after 12th grade:

EQUATION 4.1 – MULTI-DIMENSIONAL GLOBAL NET PRESENT VALUE

$$\text{NPV}(s, p) = \sum_{t=1}^T \frac{\hat{S}(s, p, t) \cdot (1 - \tau(y_t, \kappa)) \cdot P_{\text{visa}}(t, d) \cdot \text{FX}(t, \text{curr}) \cdot (1 - P_{\text{rec}}(t) \cdot \chi_{\text{loss}}) \cdot (1 - \delta_{\text{AI}}(p, t))}{\prod_{k=1}^t (1 + r(k))} - C_{\text{invest}}(p, s) - \text{OC}_{\text{PLFS}}(e, \sigma)$$

Understanding Every Part of This Equation in Plain Terms:

- $\hat{S}(s, p, t)$: **Expected Gross Salary.** Your realistic predicted salary in year t (from year 1 to 20), calculated by looking at thousands of past graduates across the entire salary spectrum—from 10th percentile to 90th percentile—rather than just the top 1% placed on the brochure.
- $\tau(y_t, \kappa)$: **Income Taxes.** What the government takes away. In India, it uses the official New Tax Regime slabs (escalating to 30% plus cess); in the US, federal, state (e.g. 0% in Texas vs 13.3% in California), and FICA taxes; in Germany, the progressive income tax.
- $P_{\text{visa}}(t, d)$: **Visa & Work Authorization Odds.** For domestic Indian colleges, this is 100% (1.0). For foreign degrees, it calculates the realistic probability of winning the US H-1B visa lottery, securing a UK Graduate Route extension, or getting an EU Blue Card.
- $\text{FX}(t, \text{curr})$: **Currency Swings (Exchange Rate).** If you study abroad and earn dollars, pounds, or euros, this models how changing currency rates affect what your savings are worth when sending money home to India.
- $P_{\text{rec}}(t) \cdot \chi_{\text{loss}}$: **Recession & Layoff Factor.** Accounts for the historical reality that an economic downturn or hiring freeze occurs roughly once every 10–12 years, causing temporary pay cuts or job search delays.
- $\delta_{\text{AI}}(p, t)$: **AI Automation Drag.** The gradual loss of salary bargaining power as automated AI tools take over routine job tasks.
- $r(k)$: **Time Value of Money (Discount Rate).** A rupee earned 15 years from now buys less than a rupee today because of inflation and interest rates.
- $C_{\text{invest}}(p, s)$: **Total True Cost of College.** Everything you pay out of pocket: tuition fees, hostel, food, coaching, textbooks, laptop, and travel across all study years.
- $\text{OC}_{\text{PLFS}}(e, \sigma)$: **Opportunity Cost.** The baseline income you could have earned if you started working right away after Class 12 instead of spending 4 years in college.

ENGINEERING INTEGRATION GUIDE • NET PRESENT VALUE ENGINE

Input Pipeline: Orchestrated within the core econometric service (`services/roi_engine.py`). Fetches wage percentiles $\hat{S}$ from `salary_trajectories`, macroeconomic rates $r(t)$ from `macro_indicators`, and tuition parameters from `cost_data`.

Execution Runtime: Vectorized NumPy operations evaluating 20-year cash flow arrays in < 5ms per program. Executed synchronously during report generation and cached in Redis with a 7-day TTL.

Output Destination: Persisted to `roi_scores.financial_roi_pct` and injected into `student_reports.results_data` for client-side rendering.

4.2 Internal Rate of Return (IRR) Formulation

The Internal Rate of Return represents the annual compounding growth rate of the money you invest in your education:

💡 IN PLAIN ENGLISH • WHAT IRR MEANS FOR PARENTS

Think of college like a bank fixed deposit or mutual fund. If an education loan charges 11% annual interest, but your degree only delivers an IRR of 7%, **you are losing money every single month**. IndiaLens warns you whenever a degree's return is lower than your education loan interest rate.

EQUATION 4.2 – INTERNAL RATE OF RETURN VIA NEWTON-RAPHSON ITERATION

$$r^* \text{ such that } \text{NPV}(s, p \mid r^*) = 0$$
$$r_{n+1}^* = r_n^* - \frac{\text{NPV}(r_n^*)}{\frac{\partial \text{NPV}}{\partial r} \big|_{r_n^*}} = r_n^* + \frac{\sum_{t=1}^T \frac{\mathcal{F}(t)}{(1+r_n^*)^t} - \mathcal{C}_{\text{total}}}{\sum_{t=1}^T \frac{t \cdot \mathcal{F}(t)}{(1+r_n^*)^{t+1}}}$$

Where $\mathcal{F}(t)$ denotes your net take-home cash flow in year t , and $\mathcal{C}_{\text{total}}$ is your total college cost plus forgone earnings. If $r^* < r_{\text{loan}}$ (your education loan interest rate), the degree destroys family wealth.

ENGINEERING INTEGRATION GUIDE • INTERNAL RATE OF RETURN

Input Pipeline: Receives the net annual cash flow array $\mathcal{F}(t)$ generated by the NPV engine. Initialized with initial guess $r_0 = 0.10$.

Execution Runtime: Solved numerically via `scipy.optimize.newton` or a custom C-accelerated loop in < 2ms, guaranteed convergence within 20 iterations.

Output Destination: Rendered as the "Yield on Education" benchmark badge on programmatic deep-dive pages.

4.3 Student-Conditioned Composite ROI Index

To provide a single 0–100 overall score while accounting for what matters personally to each student, IndiaLens adapts the weighting to your specific personality and goals:

💡 IN PLAIN ENGLISH • PERSONALIZED 0 TO 100 SCORE

Different students care about different things. An ambitious student looking to maximize cash needs a high financial weighting; a student who values work-life balance or government job stability cares more about job safety and happiness. IndiaLens customizes the 0–100 score to match **your values**, not a generic national average.

EQUATION 4.3 – STUDENT-ADAPTIVE COMPOSITE ROI INDEX

$$\Omega(s, p) = \left[\frac{\sum_{k=1}^6 w_k(s) \cdot f_k(s, p) \cdot CF_k}{\sum_{k=1}^6 w_k(s)} \right] \times 100$$

Where the weights $w_k(s)$ dynamically adapt to your answers in our adaptive quiz:

- $w_{\text{financial}}(s)$: Scales up for students focused heavily on maximizing lifetime earnings.
- $w_{\text{optionality}}(s)$: Scales up for students who want flexible careers and startup agility.
- $w_{\text{mobility}}(s)$: Rewards degrees with easy international work mobility.
- $w_{\text{safety}}(s)$: Increases stability weighting for students seeking stress-free job security.
- $w_{\text{satisfaction}}(s)$, $w_{\text{network}}(s)$: Baseline campus happiness and alumni network quality.

And the six pillar scores $f_k(s, p) \in [0, 1]$ measure:

- f_1 : Financial Return Efficiency (Net cash return relative to money invested).
- f_2 : Career Optionality Ceiling (How many high-paying career branches this degree opens).
- f_3 : Geographic & Remote Mobility (Ease of working in major metros, abroad, or remotely).
- f_4 : AI Job Resilience (Score derived from our 8-vector AI threat evaluation).
- f_5 : Student Quality of Life (Campus culture, mental health, and workload balance).
- f_6 : Alumni Leadership Density (Presence of college alumni in top executive and engineering roles).
- CF_k : Confidence Factor (Automatically discounts scores if sample sizes are small).

ENGINEERING INTEGRATION GUIDE • COMPOSITE ROI ENGINE

Input Pipeline: Dynamic weight tensor constructed from the student's 4-dimensional IRT trait vector in `pers`
`onal_intelligence.intelligence`. Component metrics $f_1 \dots f_6$ are normalized to standard normal CDF scales.

Execution Runtime: Computed in < 2ms per program. Scalable to bulk screening of 5,000+ candidate programs in < 120ms for screener views.

Output Destination: Persisted to `roi_scores.composite_score` and rendered as the primary animated radial circular gauge on dashboards.

4.4 10,000-Path Monte Carlo Debt Solvency & Default Simulation

To eliminate the nightmare of unpayable education loans, IndiaLens runs 10,000 simulated life scenarios for every student:

💡 IN PLAIN ENGLISH • TESTING YOUR LOAN AGAINST 10,000 SIMULATED FUTURES

Instead of assuming everything in life goes perfectly, our computer simulates **10,000 different possible life paths** for you. What if there is an economic recession the year you graduate? What if bank loan interest rates go up by 2%? What if you take 4 extra months to find a job? By testing your loan against 10,000 randomized futures, IndiaLens calculates the exact percentage chance your family will experience loan repayment distress—and warns you before you sign any bank loan papers.

EQUATION 4.4 – STOCHASTIC LOAN STRESS & DEFAULT PROBABILITY ENGINE

$$\begin{aligned} \text{EMI}_i(t) &= \frac{P_{\text{principal}} \cdot r_m(i, t) \cdot (1 + r_m(i, t))^n}{(1 + r_m(i, t))^n - 1} \cdot (1 + \xi_{\text{MCLR}}(i, t)) \\ d\xi_{\text{MCLR}} &= \kappa (\bar{\xi} - \xi_{\text{MCLR}}) dt + \sigma_{\text{MCLR}} dW_t \quad (\text{Floating Bank Rate Drift}) \\ \hat{S}_i(t) &= \hat{S}(t) \cdot (1 + \varepsilon_{\text{vol},i}) \cdot (1 - \delta_{\text{AI}}(t)) \cdot \mathbf{1}_{\{\text{NoRecession}_i(t)\}} + \hat{S}(t) \cdot (1 - \chi_{\text{loss}}) \cdot \mathbf{1}_{\{\text{Recession}_i(t)\}} \\ \text{DefaultEvent}_i &= \mathbf{1} \left[\exists t \in [1, n] : \text{EMI}_i(t) > \omega_{\text{distress}} \cdot \frac{\hat{S}_i(t)}{12} \cdot (1 - \tau_{\text{eff}}) \cdot (1 - r_{\text{PF}}) \right] \\ \text{Pr}[\text{Default}] &= \left(\frac{1}{N} \sum_{i=1}^{N=10,000} \text{DefaultEvent}_i \right) \times 100\% \\ T_{\text{BEP}}^{50} &= \text{median} \left\{ \min \left\{ t \geq 1 : \sum_{k=1}^t \left(\frac{\hat{S}_i(k)}{12} \cdot s_{\text{rate}} - \text{EMI}_i(k) \right) \geq \mathcal{C}_{\text{invest}}(p, s) \right\} \right\} \end{aligned}$$

Key Elements of the Simulation:

- DefaultEvent_i : Triggered whenever your monthly loan EMI exceeds 38% of your net monthly take-home salary (after taxes and mandatory Provident Fund deductions).
- $\text{Pr}[\text{Default}]$: The proportion of the 10,000 simulated life paths where the student faces loan distress. Below 5% is Safe; 5%–15% is Moderate Risk; above 20% triggers a High Risk warning.
- T_{BEP}^{50} : Median Breakeven Period—the exact number of months it will take for your accumulated net savings to fully repay your college education costs.

ENGINEERING INTEGRATION GUIDE • MONTE CARLO DEBT SIMULATOR

Input Pipeline: Ingests principal loan amount, loan tenure, and nominal interest rate from the Wizard Step 2 payload. Pulls sector volatility σ_{field} from `risk_indicators`.

Execution Runtime: Vectorized NumPy computation matrix ($10,000 \times 240$ months) executed via async FastAPI endpoint in < 65ms.

Output Destination: Yields Default Probability percentage and $P_{10}/P_{50}/P_{90}$ breakeven months rendered as an interactive solvency barometer on the student report.

SECTION 05 • AI DISRUPTION & LABOR MACROECONOMICS

The Flagship AI Risk & Future Labor Demand Index

The major economic challenge of the coming decade is how AI models, autonomous coding agents, and robotics are reshaping jobs. Traditional college counseling falsely assumes that any STEM degree will remain safe forever. IndiaLens continuously monitors real-world automation speeds at the daily job-task level to protect students from entering careers that will be automated before they graduate.

💡 IN PLAIN ENGLISH • HOW AI WILL AFFECT YOUR CAREER

AI does not usually destroy an entire profession overnight—it replaces **routine daily tasks**. For example, AI can already write basic code, translate documents, and fill out tax forms in seconds. If a job consists of 80% repetitive tasks, salaries for entry-level workers in that field will decline sharply. But if a job requires complex

human judgment, physical dexterity, negotiation, or creative system design, AI will make the human worker *more productive and higher-paid*. IndiaLens shows you which side of that line your chosen degree sits on.

5.1 Task-Level S-Curve Automation Decay Engine

IndiaLens models automation using an empirical S-curve diffusion formula that tracks how quickly companies adopt AI tools for specific tasks:

EQUATION 5.1 – TASK AUTOMATION S-CURVE DECAY FUNCTION

$$\delta_{AI}(p, t) = D_0(p) \cdot (1 - H(p)) \cdot (1 - \exp(-\lambda_{AI} \cdot t)) \cdot \mathcal{M}_{adopt}(p, t)$$
$$\mathcal{M}_{adopt}(p, t) = \frac{1}{1 + \exp(-k_p(t - t_{0,p}))} \quad (\text{Sector-Specific Adoption Envelope})$$

Understanding the Terms in Plain English:

- $D_0(p)$: **Task Exposure Baseline.** What percentage of this job involves routine cognitive work that modern AI can do (e.g., 75% for routine software testing vs. 15% for pediatric surgery).
- $H(p)$: **The Human Moat.** The share of tasks requiring deep human empathy, physical presence, ethics, or legal accountability that robots and software cannot replace.
- λ_{AI} : **AI Improvement Speed.** The measured annual rate at which AI models become faster, cheaper, and smarter.
- $\mathcal{M}_{adopt}(p, t)$: **Real-World Corporate Adoption.** How fast actual companies in that industry are buying and implementing AI tools (e.g., fast in customer service; much slower in licensed civil engineering).

5.2 The Eight-Vector Labor Risk Surface

To give students a complete picture of job stability, IndiaLens tracks eight separate risk factors for every profession through automated weekly data feeds:

TABLE 5.1 – THE EIGHT ORTHOGONAL RISK VECTORS OF THE AI RISK SURFACE

Risk Vector (V_i)	What It Measures in Everyday Terms	Where Data Comes From	Calculation Method
1. AI Automation (V_1)	Fraction of daily tasks that can be performed by current AI models	Oxford O*NET, ILO Research & ArXiv	w_1^R (Ridge Regression)
2. Salary Volatility (V_2)	How wildly salaries swing from good years to bad years	AmbitionBox 5-Year Salary Data	w_2^R (Ridge Regression)
3. Industry Cyclicalilty (V_3)	Vulnerability to broader economic recessions and market downturns	Reserve Bank of India (RBI) Database	w_3^R (Ridge Regression)
4. Degree Oversupply (V_4)	Whether too many colleges are pumping out graduates compared to actual jobs	Ministry of Education AISHE vs. Naukri	w_4^R (Ridge Regression)
5. Location Concentration (V_5)	Whether all high-paying jobs are stuck in just 2 cities or spread across India	Naukri, Indeed & LinkedIn Insights	w_5^R (Ridge Regression)
6. Government & Visa Policy (V_6)	Risk of sudden government rule changes, quota shifts, or foreign visa caps	Gazette of India, NMC Regulations, US FR	w_6^R (Ridge Regression)
7. Physical Burnout (V_7)	Physical fatigue, hazardous working conditions, and health strain	Ministry of Labour Safety Records	w_7^R (Ridge Regression)
8. Work-Life Balance (V_8)	Average working hours, toxic work cultures, and employee burnout rates	AmbitionBox, Glassdoor & Reddit	w_8^R (Ridge Regression)

EQUATION 5.2 – RIDGE-LEARNED COMPOSITE RISK TENSOR

$$\mathcal{R}(p) = \sum_{i=1}^8 w_i^{\mathcal{R}} \cdot V_i(p) \quad \text{where} \quad \mathbf{w}^{\mathcal{R}} = \arg \min_{\mathbf{w}} \|\mathbf{V}\mathbf{w} - \mathbf{y}_{\text{instability}}\|_2^2 + \lambda_{\text{ridge}} \|\mathbf{w}\|_2^2$$

Weights are re-calibrated weekly against real salary histories of alumni cohorts to ensure the risk score reflects actual job market stability rather than subjective guesswork.

5.3 The Job Security Score (JSS) & Tail Layoff Probabilities

To convert these complex risk factors into an actionable 0–100 score that anyone can understand instantly, IndiaLens produces the **Job Security Score (JSS)**:

💡 IN PLAIN ENGLISH • READING YOUR JOB SECURITY SCORE

The Job Security Score ranges from **35 (High Risk)** to **99 (Maximum Security)**:

- **Score 85–99 (Green):** Extremely secure career with strong human moats (e.g., specialized medical surgery, high-level system architecture, civil infrastructure). AI acts as an assistant that boosts your income.
- **Score 65–84 (Yellow):** Moderate automation. You will need occasional upskilling, but the field offers steady employment.
- **Score Below 65 (Red Alert):** High automation vulnerability. Over half of entry-level tasks can already be done by software. IndiaLens warns you and shows the exact emergency savings you must keep for mid-career retraining.

EQUATION 5.3 – THE JOB SECURITY SCORE & ACTUARIAL LAYOFF HORIZON

$$\text{JSS}(s, p) = \text{clip}_{35}^{99} [\sigma(A \cdot \Theta_{\text{res}}(p) + B \cdot \theta_{\text{ai}}(s) + C \cdot \text{Tier}_{\text{boost}}(p) - D \cdot \delta_{\text{AI}}(p, 0)) \times 100]$$

$$\Theta_{\text{res}}(p) = w_{\phi} \cdot \Phi_{\text{phys}}(p) + w_{\gamma} \cdot \Gamma_{\text{cog}}(p) + w_{\eta} \cdot \text{H}_{\text{soc}}(p) + w_{\rho} \cdot \text{P}_{\text{reg}}(p)$$

$$\hat{\mathcal{L}}_{5Y}(p, s) = \max(1.0, \min(35.0, (100 - \text{JSS}) \cdot 0.35 \cdot (1 + 0.5 \cdot \sigma_{\text{wage}})))$$

$$\hat{\mathcal{L}}_{10Y}(p, s) = \max(2.5, \min(50.0, (100 - \text{JSS}) \cdot 0.60 \cdot (1 + \mathcal{M}_{\text{adopt}}(p, 10))))$$

$$\text{UpskillingReserve}_{\text{INR}}(p, s) = ₹1,20,000 \times \left(1 + \frac{100 - \text{JSS}}{100}\right) \times (1 + \pi_{\text{ed}})^5$$

Where the Resilience Tensor Θ_{res} isolates four essential human advantages: Physical Task Dexterity (Φ_{phys}), Complex Non-Routine Problem Solving (Γ_{cog}), Social & Negotiation Empathy (H_{soc}), and Legal Regulatory Licensing (P_{reg}).

ENGINEERING INTEGRATION GUIDE • AI RISK & JSS ENGINE

Input Pipeline: Pulls 8 orthogonal vectors from `risk_indicators` table. For novel unindexed job titles, executes zero-shot semantic projection via `POST /api/v2/ai/evaluate-career-string` in < 800ms.

Execution Runtime: Computed in batch every Sunday via Airflow DAG and on-demand for customized student pathways.

Output Destination: Persisted to `risk_indicators.jss_score` and rendered as an interactive 8-Axis Radar Chart on report screens.

5.4 Zero-Hardcoding Mandate & Fluid Real-Time Recalibration

A fundamental rule of IndiaLens's architecture is that nothing is hardcoded or based on static tables. Everything updates dynamically:

- **Instant Analysis for New Job Titles:** If a student enters a brand-new or unusual career (like *"Quantum Machine Learning Specialist"* or *"AI Bio-Robotics Engineer"*), IndiaLens automatically scans recent research papers, job listings, and patent filings to calculate its 8-vector risk score in under 800 milliseconds.
- **Weekly Automated Updates:** Macroeconomic data (bank interest rates, inflation figures) and job openings from LinkedIn and Naukri are refreshed automatically every Sunday night, ensuring students always see the current job market reality.

SECTION 06 • PSYCHOMETRIC ENGINE

The 3PL IRT Computerized Adaptive Testing (CAT) Engine

Traditional aptitude tests make students sit through boring 100-question multiple choice quizzes that cause mental exhaustion and encourage lucky guesses. **IndiaLens** replaces this with an intelligent Computerized Adaptive Testing (CAT) engine powered by 3-Parameter Logistic (3PL) Item Response Theory (IRT). It measures four essential student strengths in just 6 to 8 interactive micro-dilemma scenarios.

💡 IN PLAIN ENGLISH • WHY OUR QUIZ DISCOVERS REAL TALENT IN 8 QUESTIONS

Think of this quiz like playing chess against a grandmaster computer. When you answer a question correctly, the computer doesn't waste your time with more easy questions—it immediately gives you a slightly more challenging dilemma to test the limits of your abilities. If you get stuck or pick a cautious answer, it gives you an easier scenario to find your baseline. Instead of asking 50 generic questions, **in just 8 smart questions** IndiaLens pinpoints your true problem-solving stamina, risk tolerance, economic motivation, and tech agility—while filtering out lucky random clicks.

6.1 The 3-Parameter Logistic (3PL) Mathematical Model

The probability that a student j with underlying strength level θ_j selects the optimal answer on question i is modeled mathematically as:

EQUATION 6.1 – 3PL ITEM RESPONSE PROBABILITY FUNCTION

$$P(x_{ij} = 1 \mid \theta_j, a_i, b_i, c_i) = c_i + \frac{1 - c_i}{1 + \exp(-a_i(\theta_j - b_i))}$$

Understanding the Three Parameters in Simple Terms:

- $a_i > 0$: **Discrimination Power.** How sharply this specific question separates students with high aptitude from students who are struggling. High- a questions give maximum diagnostic insight.
- $b_i \in [-3.0, +3.0]$: **Question Difficulty.** How tough the question is. An average student has a 50% chance of solving a question with difficulty b_i .
- $c_i \in [0, 0.25]$: **Lucky Guessing Protection.** Accounts for students who randomly pick an answer or make a careless slip, ensuring a single mistake does not unfairly hurt their overall score.

6.2 Fisher Information Maximization & Dynamic Item Selection

Rather than giving every student the exact same fixed set of questions, the test dynamically selects whichever next question provides the highest diagnostic certainty:

EQUATION 6.2 – FISHER INFORMATION & OPTIMAL ITEM SELECTION

$$\mathbf{I}_i(\boldsymbol{\theta}) = a_i^2 \cdot \frac{P_i(\boldsymbol{\theta}) - c_i}{(1 - c_i)^2} \cdot \frac{(1 - P_i(\boldsymbol{\theta}))}{P_i(\boldsymbol{\theta})} \cdot \mathbf{u}_i \mathbf{u}_i^T$$

$$\text{Item}_{t+1} = \arg \max_{i \in \text{Unadministered}} \det \left(\sum_{k \in \text{Administered}} \mathbf{I}_k(\hat{\boldsymbol{\theta}}_t) + \mathbf{I}_i(\hat{\boldsymbol{\theta}}_t) \right)$$

This mathematical selection rule guarantees that every single question presented extracts the maximum possible insight about your natural instincts and career preferences.

6.3 Expected A-Posteriori (EAP) Trait Updates

After each answer you give, IndiaLens instantly updates your profile using Bayesian mathematics:

EQUATION 6.3 – GAUSS-HERMITE QUADRATURE TRAIT ESTIMATION & STOPPING RULE

$$\hat{\boldsymbol{\theta}}_{\text{EAP}} = \frac{\sum_{q=1}^{41} \mathbf{X}_q \cdot \mathcal{L}(\mathbf{x}_{1:t} | \mathbf{X}_q) \cdot \mathbf{W}_q}{\sum_{q=1}^{41} \mathcal{L}(\mathbf{x}_{1:t} | \mathbf{X}_q) \cdot \mathbf{W}_q} \quad \text{where} \quad \mathcal{L}(\mathbf{x}_{1:t} | \mathbf{X}_q) = \prod_{k=1}^t P_k(\mathbf{X}_q)^{x_k} (1 - P_k(\mathbf{X}_q))^{1-x_k}$$

$$\text{Standard Error: } \text{SE}(\hat{\boldsymbol{\theta}}_d) = \frac{1}{\sqrt{\sum_{k=1}^t \mathbf{I}_{k,d}(\hat{\boldsymbol{\theta}})}}$$

$$\text{Diagnostic Termination Criterion: } \max_{d \in \{1,2,3,4\}} \text{SE}(\hat{\boldsymbol{\theta}}_d) < 0.35 \quad \text{OR} \quad n_{\text{items}} \geq 8$$

The Four Core Student Traits Measured (on a -3.0 to +3.0 scale):

1. θ_{risk} : **Risk Tolerance.** Are you comfortable with startups and performance bonuses, or do you prefer the steady security of government/corporate roles?
2. θ_{value} : **Financial Motivation.** How strongly do you prioritize high salaries and fast loan repayment versus lifestyle, work-life balance, and personal passion?
3. θ_{autonomy} : **Independence vs. Structure.** Do you thrive in self-directed entrepreneurial environments or structured corporate hierarchies?
4. θ_{ai} : **AI & Tech Adaptability.** How quickly and comfortably do you adopt new AI tools and technical workflows?

ENGINEERING INTEGRATION GUIDE · 3PL IRT ADAPTIVE TESTING MODULE

Input Pipeline: Step 3 of the Wizard calls `POST /api/v2/psychometrics/next-item` with the current session token and response binary $x_t \in \{0,1\}$.

Execution Runtime: Gauss-Hermite integration executes in < 15ms. Evaluates stopping criteria $\text{SE} < 0.35$ or $n \geq 8$; if complete, marks `is_complete: true` and persists final 4-trait vector.

Output Destination: Persisted to `personal_intelligence.intelligence` JSONB; dynamically modulates all downstream Composite ROI weights $w_k(s)$.

Machine Learning Model Specifications & Pipelines

IndiaLens uses production-grade machine learning models to turn millions of verified career data points into accurate, personalized predictions. Rather than guessing a single salary number, our models calculate your full 20-year career earnings curve across different scenarios.

💡 IN PLAIN ENGLISH · HOW WE PREDICT REAL CAREER SALARIES

When a college brochure says *"Average Package ₹12 LPA"*, they often conceal that 50% of the batch got ₹4 LPA while two students who got overseas jobs artificially inflated the average. IndiaLens trains machine learning algorithms on thousands of verified alumni histories. We look at your specific branch, college tier, city, and economic conditions to predict what you will realistically make at Year 1, Year 3, Year 5, Year 10, and Year 20—showing the full spread from cautious cases (10th percentile) to outstanding outcomes (90th percentile).

7.1 Multi-Horizon XGBoost Quantile Salary Regressor

IndiaLens deploys gradient-boosted decision trees trained with the Pinball (Quantile) Loss function across five percentile levels ($P_{10}, P_{25}, P_{50}, P_{75}, P_{90}$) over 8 key career milestones (Year 1 to Year 20):

EQUATION 7.1 – QUANTILE (PINBALL) LOSS FORMULATION

$$\mathcal{L}_\alpha(y, \hat{y}_\alpha) = \max(\alpha(y - \hat{y}_\alpha), (\alpha - 1)(y - \hat{y}_\alpha))$$

Hyperparameter Grid & Training Configuration:

- `max_depth: 6`, `learning_rate: 0.05`, `n_estimators: 450`, `min_child_weight: 5`
- `subsample: 0.85`, `colsample_bytree: 0.80`, `reg_lambda (L2): 1.8`, `reg_alpha (L1): 0.4`
- Input Feature Matrix (18 Verified Indicators):** `degree_field_id`, `institutional_tier` (1/2/3), `college_type` (IIT/NIT/BITS/State/Private), `nirf_score`, `naac_grade_num`, `established_year`, `placement_rate_pct`, `total_cost_inr`, `ai_automation_prob`, `salary_volatility`, `industry_gva_growth`, `city_tier`, `alumni_leadership_density`, `internship_rate`, `faculty_student_ratio`, `h_index_citations`, `state_gdp_per_capita`, `macro_cpi_rate`.

SECTION 08 · DUAL-ENGINE AI CORE

The Student Intelligence Layer: Dual-Engine Architecture (DeepSeek + Sonar)

The Student Intelligence Layer represents IndiaLens's conversational and reasoning pinnacle. Generic AI chatbots fail in high-stakes educational advisory because they act as sycophants—hallucinating fake admissions cutoffs, flattering user biases, and producing agreeable cheerleading without mathematical grounding. Monolithic frontier models (such as Gemini 3.7 Flash with Google Search Grounding) introduce severe cost bloat (\$35 per 1,000 search queries on top of token fees) and combine memory recall, web search, and numerical reasoning into a single fragile bottleneck. IndiaLens solves this by deploying a **decoupled, dual-engine ML-LLM architecture** pairing **DeepSeek V4 Flash** (Personal Mentor & Memory Core) with **Perplexity Sonar** (Live Web Detective), orchestrated through a high-performance **Fork-and-Join pipeline**.

💡 IN PLAIN ENGLISH • HOW OUR TWO AI ENGINES WORK TOGETHER

Think of this architecture like a team of two trusted advisors:

- **DeepSeek V4 Flash is your personal mentor:** It remembers your family's budget, your test scores, your strengths, and your past questions by heart.
- **Perplexity Sonar is the online detective:** It searches the live internet in real time, pulling verified circulars from college websites, checking actual cutoff ranks, and uncovering campus realities.
- **The Smart Handshake:** Sonar sends all hard numbers (tuition fees, closing ranks) directly to our Python mathematical calculator, while sending qualitative campus tips (like strict attendance rules or hostel shortages) directly to DeepSeek. DeepSeek then synthesizes everything into a clear, honest, and personal report—telling you the truth without sugarcoating.

8.1 The Dual-Engine Intelligence Hierarchy

Instead of overloading one AI model with every task, IndiaLens cleanly divides work across three specialized layers:

TABLE 8.1 – DECOUPLED DUAL-ENGINE INTELLIGENCE SPECIFICATION

Intelligence Layer	Underlying Technology	Role & Core Responsibility in Plain English
Layer 1: Deterministic Math Calculator	NumPy, SciPy, XGBoost, Monte Carlo Engine	Calculates exact salary curves, loan default risks, and cutoff scores in < 35ms. AI language models are strictly forbidden from doing math calculations to prevent hallucinations.
Layer 2A: The Web Detective	Perplexity Sonar Search API (Live Citations & Web Index)	Carries out live, zero-hallucination web research across official college portals (JoSAA, JAC, university circulars, visa guidelines). Retrieves up-to-date fee hikes and real closing ranks with verified website links.
Layer 2B: The Personal Mentor	DeepSeek V4 Flash (128k+ Context, Fast Memory Retrieval)	Remembers your background and family budget, directs Sonar on what to search for (Pass 1), and combines the math calculator's output with campus tips to write you a clear, honest advisory report (Pass 2).

8.2 The 5-Step Fork-and-Join Pipeline

To eliminate delays and ensure maximum accuracy, data moves in 5 clear, unidirectional steps:

- STEP 1

Understanding Your Question (DeepSeek Pass 1): You ask a question. DeepSeek looks up your personal profile, test scores, and budget limits, converting your question into specific factual search terms for Sonar.
- STEP 2

Searching Live Verified Sources (Perplexity Sonar): Sonar searches verified official websites and extracts accurate, up-to-date numbers along with direct source links.
- STEP 3

The Smart Fork (Splitting Numbers and Tips):
 - STREAM A (HARD NUMBERS ONLY)** : Clean numbers (tuition fees, hostel costs, recent cutoff ranks) are sent straight to our Python math calculator.
 - STREAM B (CAMPUS TIPS & CONTEXT)** : Real-world ground facts (e.g. strict 75% attendance rules, placement drive changes, early scholarship deadlines) bypass the calculator and go directly to DeepSeek.
- STEP 4

Running the Math (Python Actuarial Engine): In under 35 milliseconds, the math calculator computes your 20-year net return, solves the loan breakeven time, and tests your loan against 10,000 simulated life paths.
- STEP 5

The Honest Synthesis (DeepSeek Pass 2): DeepSeek takes the math results, adds the campus tips, and writes a clear, personalized report highlighting financial risks first, followed by growth opportunities.

8.3 Multi-Stream Data Contracts (The System Handshake)

The system prevents AI hallucinations by strictly separating numbers from narrative text:

```
// Stream A: Clean Numbers (Sonar → Python Math Calculator)
STREAM_A_PARAMETRIC = {
  "program_id": "THAPAR-COE-CS",
  "tuition_annual_inr": 540000.0,
  "hostel_annual_inr": 120000.0,
  "duration_years": 4.0,
  "closing_rank_recent": 16500,
  "verified_median_salary_inr": 1150000.0
}

// Stream B: Real-World Campus Facts (Sonar → DeepSeek Direct Bypass)
STREAM_B_QUALITATIVE = {
  "hiring_alerts": [
    "IT mass recruiters reduced on-campus intake by 40% in recent cycle.",
    "Strict 75% attendance policy enforced; low flexibility for off-campus internships."
  ],
  "tactical_tips": [
    "Apply in Round 1 counseling before May 15 to lock the 20% board merit waiver.",
    "Verify campus location: Patiala main campus vs Derabassi satellite center."
  ],
  "source_citations": [
    "https://thapar.edu/admissions/circular-2025.pdf",
    "https://reddit.com/r/ThaparUniversity"
  ]
}
```

```
// Immutable Context Packet Injected into DeepSeek Pass 2
CONTEXT_PACKET = {
  "student_token": "uuid-v4-9b1deb4d-3b7d-4bad-9bdd-2b0d7b3dcb6d",
  "psychometric_profile": {
    "theta_risk": -1.42,      // Highly risk-averse; requires strong safety margins
    "theta_value": 2.15,     // Strong economic maximizer; prioritizes NPV & short breakeven
    "theta_autonomy": 0.35,  // Moderate corporate preference
    "theta_ai": 1.80        // High technological adaptability & prompt literacy
  },
  "financial_constraints": {
    "max_household_budget_inr": 1800000,
    "max_acceptable_loan_inr": 1200000,
    "max_monthly_emi_tolerance_inr": 28000,
    "liquidity_buffer_months": 6
  },
  "actuarial_ml_outputs": {
    "composite_roi_score": 88.4,
    "npv_20y_inr": 18450000,
    "irr_pct": 21.4,
    "monte_carlo_default_prob_pct": 1.2,
    "job_security_score": 82.0,
    "admission_tier": "Target (68% admission probability)"
  },
  "stream_b_qualitative": STREAM_B_QUALITATIVE
}
```

8.4 Anti-Yes-Man Fiduciary Directives (No Sugarcoating Allowed)

To ensure students receive honest, uncompromised advice, DeepSeek's prompt is locked with strict rules:

CORE SYSTEM PROMPT GUARDRAILS (STRICT ADVISORY RULES)

- **The Student First:** Your sole duty is to protect the student's 20-year financial health and career resilience. You are strictly forbidden from flattering the student, agreeing with unrealistic plans, or cheerleading bad decisions.
- **Risks First, Rewards Second:** Whenever evaluating a degree, you must explain the worst-case scenario, the loan EMI burden, and the risk of job market struggles BEFORE discussing the upside.
- **Calling Out Dangerous Decisions:** If a cautious student with limited family savings plans to borrow ₹60 Lakhs for an overseas degree with high visa rejection rates, you must intervene and explicitly calculate their danger of bankruptcy.
- **Zero Fabricated Numbers:** You may never make up salary, placement, or cutoff figures. All financial data must come strictly from the math calculator or verified source links.

8.5 The Dynamic Path Graph (Longitudinal Career Memory)

IndiaLens tracks your entire academic and career journey as a dynamic roadmap in our database:

- **Roadmap Milestones:** Records exam target milestones (JEE, SAT, IELTS), enrolled degrees, acquired technical skills (PyTorch, cloud architecture), and your 10-year target job role.
- **Automatic Recalibration:** When you update your mock exam score or family budget, or when industry hiring trends shift, IndiaLens automatically adjusts your roadmap, pruning risky options and highlighting better alternatives.

8.6 Stack Comparison: DeepSeek + Sonar vs. Legacy Models

Comparing our dual-engine setup against traditional single models demonstrates massive cost savings and superior reliability:

TABLE 8.2 – OPERATIONAL & FINANCIAL STACK COMPARISON

Performance Metric	Legacy Single Model (e.g. Gemini 3.7 Flash)	IndiaLens Production Stack (DeepSeek + Sonar)
Cost per 1,000 Advisory Chats	\$38.00 – \$55.00 (Expensive search surcharge)	\$2.80 – \$4.20 (92% Cost Savings)
Information Freshness	Generic web search wrapper	Dedicated Perplexity Sonar search index with verified URLs
Memory & Conversation Cost	Gets progressively expensive as chat grows	Ultra-low caching costs (< \$0.02 per 1M cached tokens)
Response Speed	1,400ms – 2,200ms (Slow combined reasoning)	< 950ms Cold / < 420ms Warm (Instant response)
Math Accuracy	Risk of language model hallucinating numbers	100% Accurate (Python calculator handles all math)

SECTION 09 · ADMISSIONS PORTFOLIO ENGINEERING

The Caliber-Calibrated Admissions Portfolio Engine

A fatal mistake in traditional counseling is giving students unrealistic, generic advice (like telling a student scoring in the 80th percentile to "just work harder and crack IIT Bombay Computer Science"). **IndiaLens** deploys an empirical **Caliber-Calibrated Admissions Portfolio Engine** that maps your actual test scores and grades into a balanced, 4-tier college application strategy that protects you from heartbreak.

💡 IN PLAIN ENGLISH · WHY YOU NEED A 4-TIER COLLEGE PORTFOLIO

Think of college applications like building a smart investment portfolio. If you only apply to top-10 "dream" colleges, a single bad exam day could leave you with zero college seats. If you only apply to easy colleges, you sell yourself short. IndiaLens divides your college choices into 4 clear categories: (1) **Reach Colleges** (your dream schools), (2) **Target Colleges** (schools where your current scores give you a strong 60%–75% chance), (3) **Safety Colleges** (guaranteed admission even if exam day goes poorly), and (4) **Hidden Gems** (lesser-known colleges with great placements and low tuition that beat expensive private universities).

9.1 The Cutoff Delta Function & Z-Score Formulation

Admission probabilities are computed by comparing historical cutoff ranks with your expected exam performance:

EQUATION 9.1 – CUTOFF DELTA Z-SCORE & ADMISSION PROBABILITY

$$Z(s,p) = \frac{\text{HistoricalClosingRank}(p, \text{cat}, \text{quota}) - \mathbb{E}[\text{StudentRank}(s)]}{\sigma_{\text{exam_variability}}}$$
$$\Pr[\text{Admission}(s,p)] = \Phi(Z(s,p)) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{Z(s,p)} e^{-u^2/2} du$$

Understanding the Terms in Plain English:

- **HistoricalClosingRank:** The 3-year rolling average closing rank for that college branch under your category reservation (General, OBC, SC, ST, EWS) and home-state quota.
- **$\mathbb{E}[\text{StudentRank}(s)]$:** Your predicted rank based on your latest mock tests and preparation trajectory.
- **$\sigma_{\text{exam_variability}}$:** Normal exam-day fluctuation (typically 1,500 to 2,000 ranks around the 95th percentile in JEE Main).
- **$\Pr[\text{Admission}]$:** Your exact mathematical probability of securing a seat.

ENGINEERING INTEGRATION GUIDE · ADMISSIONS Z-SCORE CALIBRATION

Input Pipeline: Receives expected rank $\mathbb{E}[\text{Rank}]$ from Wizard Step 1 and queries historical 3-year closing cutoffs from `programs.cutoff_rank_high` filtered by category and home state quota.

Execution Runtime: Standard normal CDF $\Phi(Z)$ calculated in < 1ms via `scipy.stats.norm.cdf`.

Output Destination: Segregates programs into Reach, Target, Safety, and Hidden Gem tabs rendered on the student's personal portfolio view.

9.2 The Four Portfolio Tiers

Based on your calculated admission probability, colleges are organized into a diversified list:

TABLE 9.1 – THE FOUR CALIBER-CALIBRATED ADMISSIONS TIERS

Portfolio Tier	Admission Probability	Statistical Bounds	Role in Your College Strategy
1. Reach Tier	15% – 25%	$-1.5 \leq Z < -0.67$	Ambitious dream schools. Worth trying if you have an exceptional test day. Capped at 20% of your total applications.
2. Target Tier	60% – 75%	$-0.25 \leq Z \leq +0.67$	Core alignment. Matches your consistent mock exam scores. Should form 50% of your application list.
3. Safety Tier	90% – 98%	$Z > +1.28$	Downside insurance. Guaranteed admission even under severe test-day stress or illness. Forms 20% of your list.
4. Hidden Gem Arbitrage	75% – 95%	$Z > +0.67$ AND $\Omega \geq 85$	High-value public or autonomous colleges offering exceptional placement returns and low fees, free from private college marketing hype.

9.3 Quota Strategy & Smart Budget Pruning

- **State Quota Advantage:** Automatically checks home-state reservation benefits (e.g., Delhi students at DTU, NSUT, and IPU enjoying 85% regional seat quotas).
- **Branch vs. College Brand Trade-Off:** Evaluates whether taking a core engineering branch at a premier Tier-1 college (like Mechanical at IIT Kharagpur or Chemical at BITS Pilani) yields a better 20-year career return than Computer Science at an expensive Tier-3 private college.
- **Automatic Budget Protection:** Any college whose total 4-year cost exceeds 130% of your family's stated budget without qualifying for low-interest education loans is automatically flagged with a clear affordability

warning.

Global Degree ROI & Cross-Border Actuarial Models

Studying abroad is one of the largest financial investments an Indian family can make, often requiring ₹40 Lakhs to ₹1.2 Crores in savings or bank loans. Traditional study-abroad agents push colleges solely based on secret commission kickbacks. **IndiaLens** evaluates foreign degrees with strict financial math, factoring in visa lottery odds, actual city rents, income taxes, and currency exchange rates.

💡 IN PLAIN ENGLISH • EXPOSING STUDY-ABROAD AGENT SCAMS

Study abroad agents frequently push students into expensive 1-year degrees in the US or UK without explaining that: (1) **A Non-STEM US degree only gives you 1 year to work and a 25% chance of winning an H-1B visa**—meaning a 75% chance you are forced to return to India with an ₹80 Lakh dollar-debt that is impossible to repay on Indian salaries; (2) **Earning \$120k in San Francisco leaves you with LESS savings than earning \$95k in Texas or €55k in Germany** because California taxes and Silicon Valley rents consume almost your entire paycheck; (3) **Public German universities charge ZERO tuition** and offer a 94% fast-track to permanent residency, giving lower-risk returns than many mid-tier American schools.

10.1 Cumulative STEM OPT & H-1B Visa Survival Probability

In the United States, your ability to pay back a loan depends on securing an H-1B work visa before your student work permit (OPT) expires:

EQUATION 10.1 – CUMULATIVE H-1B LOTTERY SURVIVAL PROBABILITY

$$P_{H-1B}^{(\text{cumulative})} = 1 - (1 - p_{\text{lottery}})^n \quad \text{where } p_{\text{lottery}} \approx 0.25 \text{ (Annual Selection Rate)}$$

STEM Degree (3-Year OPT, $n = 3$ Lottery Attempts): $P_{H-1B}^{(3)} = 1 - (1 - 0.25)^3 = 57.8\%$

Non-STEM Degree (1-Year OPT, $n = 1$ Lottery Attempt): $P_{H-1B}^{(1)} = 1 - (1 - 0.25)^1 = 25.0\%$

A Non-STEM degree gives you only one lottery attempt (a **75% chance of being forced to leave** within 12 months), creating massive financial danger for families who took education loans.

ENGINEERING INTEGRATION GUIDE • GLOBAL VISA RISK ENGINE

Input Pipeline: Fetches `is_stem_designated` boolean and country code from `global_programs` table.

Execution Runtime: Injected as the $P_{\text{visa}}(t, d)$ discount multiplier into the master Cross-Border NPV equation.

Output Destination: Expressed as a prominent visual "Visa Survival Probability" risk metric on all overseas degree cards on `/global`.

10.2 City-Level Cost of Living & Real Annual Savings

A \$120,000 salary in an expensive metro yields far less real savings than a \$95,000 salary in a low-tax state:

EQUATION 10.2 – SPATIAL NET ANNUAL CASH FLOW & SAVINGS FORMULATION

$$\text{NetAnnualSavings}(c, t) = \hat{S}(p, c, t) \cdot (1 - \tau_{\text{total}}(c)) - 12 \cdot \text{Rent}_{\text{median}}(c) - 12 \cdot \text{COLI}_{\text{baseline}}(c)$$
$$\tau_{\text{total}}(c) = \tau_{\text{federal}}(y) + \tau_{\text{state}}(c, y) + \tau_{\text{FICA}} + \tau_{\text{local}}(c)$$

Comparing Real Metros in Plain Numbers:

- **San Francisco, CA:** Total taxes $\approx 40.5\%$, Rent $\approx \$3,200/\text{month}$, living expenses $\approx \$1,400/\text{month}$. On a $\backslash\$120\text{k}$ salary, your **net savings are only** $\approx \$16,200/\text{year}$.
- **Austin, TX:** Total taxes $\approx 23.2\%$ (Zero state income tax), Rent $\approx \$1,450/\text{month}$, living expenses $\approx \$950/\text{month}$. On a $\backslash\$100\text{k}$ salary, your **net savings are** $\approx \$48,000/\text{year}$ (**3x higher net savings**).
- **Munich, Germany:** Total taxes $\approx 34.0\%$, Rent $\approx \text{€}1,300/\text{month}$, living expenses $\approx \text{€}850/\text{month}$. Zero tuition fees + EU Blue Card fast-track to permanent residency in 21 months (94% visa security).

10.3 The Four-Tier Global Institution Taxonomy

TABLE 10.1 – GLOBAL DEGREE STRATEGIC TAXONOMY

Global Tier	Representative Universities	Tuition & Living Costs	IndiaLens Recommendation
Tier 1: Value Champions	Purdue, Georgia Tech, UIUC, Texas A&M, TU Delft	Moderate tuition ($\backslash\$28\text{k}\text{--}\backslash\$38\text{k}/\text{yr}$), high STEM placement rates, huge alumni network	Strong Buy: Outstanding return on investment with a median 3.5-year loan breakeven.
Tier 2: Zero-Tuition Public Universities	TU Munich, RWTH Aachen, Heidelberg, KTH Sweden	Zero or low tuition ($\text{€}0\text{--}\text{€}3\text{k}/\text{yr}$), living costs only, 94% EU Blue Card visa retention	Best Financial Safety: Ideal for families with budget limits; near-zero risk of loan distress.
Tier 3: Elite Brand Ceiling	Stanford, MIT, CMU, Oxford, Imperial College	High tuition ($\backslash\$60\text{k}\text{--}\backslash\$75\text{k}/\text{yr}$), extraordinary upside into global AI re-search and startups	Selective Buy: Recommended for top academic candidates aiming for high-growth tech careers.
Tier 4: Danger Zone (Diploma Mills)	Unaccredited private colleges, regional non-STEM diplomas	High tuition ($\backslash\$30\text{k}\text{--}\backslash\$45\text{k}/\text{yr}$), low placement rates, 1-year OPT, high loan default danger	Strong Avoid: Negative expected return; flagged with red alert warnings on IndiaLens.

SECTION 11 • INTERNATIONAL ADMISSIONS STRATEGY

The Global Portfolio Builder & Admissions Spike System

Top international universities (Ivy League, Stanford, Oxford, Cambridge, NUS) reject over 90% of applicants who submit generic "well-rounded" resumes (like playing an instrument, general school sports, and generic volunteering certificates). Holistic admissions officers look for an **Angular Spike**—genuine, focused excellence in one specific subject or technical area. **IndiaLens** includes an automated Portfolio Builder that helps students craft authentic, high-impact profiles.

💡 IN PLAIN ENGLISH • WHAT IS AN "ANGULAR SPIKE"?

Admissions committees read thousands of essays that all sound the same: *"I was the school prefect, I like coding, and I volunteered for an NGO for two weeks."* These generic profiles are immediately rejected. An **Angular Spike** means showing deep, continuous passion in ONE specific field. For example, instead of claiming you love biology, you build an open-source tool that analyzes local water pollution data, test it across 10 villages, and publish your findings. IndiaLens helps you plan and execute a genuine spike project from 9th grade through 12th grade.

11.1 The Angular Spike Authenticity Formulation

A student's portfolio is evaluated through an authenticity score that rewards genuine commitment over superficial resume padding:

EQUATION 11.1 – ANGULAR SPIKE AUTHENTICITY SCORE

$$\mathcal{A}_{\text{spike}}(s) = \sum_{a \in \text{Activities}(s)} \frac{1}{N_a} \cdot \mathcal{E}_a \cdot \mathcal{D}_a \cdot \cos(\mathbf{v}_a, \mathbf{u}_{\text{target_major}})$$

Understanding What the Formula Rewards:

- $\frac{1}{N_a}$: **Rarity of Activity**. Doing something unique (like building a working solar battery controller) scores much higher than common activities that 50,000 other students list.
- $\mathcal{E}_a$: **External Validation**. Winning a recognized science olympiad, getting a research paper indexed, or receiving a government grant.
- $\mathcal{D}_a$: **Multi-Year Commitment**. Dedicating 2 years of steady effort to a project scores 5x higher than cramming a project in the month before college applications.
- $\cos(\mathbf{v}_a, \mathbf{u}_{\text{target_major}})$: **Major Alignment**. How directly your activities connect to your chosen college degree.

ENGINEERING INTEGRATION GUIDE • SPIKE AUTHENTICITY ENGINE

Input Pipeline: Triggered on `/portfolio-builder` when student adds or modifies activities. Activity descriptions are embedded using `sentence-transformers` to construct $\mathbf{v}_a$ and matched against $\mathbf{u}_{\text{major}}$.

Execution Runtime: Vector similarity and formula computation executed in < 35ms.

Output Destination: Persisted to `portfolio_profiles.spike_authenticity_score` and rendered as an interactive live gauge on the Portfolio Studio.

11.2 The Action-Impact X-Y-Z Resume Standard

Every activity bullet on your resume is rewritten using the proven Google/CommonApp X-Y-Z formula:

THE X-Y-Z RESUME TRANSFORMATION STANDARD

"Accomplished [X], as measured by quantitative metric [Y], by executing technical/strategic initiative [Z]."

Weak Generic Resume Bullet: "Built a machine learning project for water quality monitoring."

IndiaLens Optimized Bullet: "Engineered an edge-deployed computer vision model running on Raspberry Pi, detecting bacterial water contamination with 94.2% precision across 14 rural wells, reducing local contamination response time by 72% in collaboration with regional health authorities."

11.3 4-Year Milestone Framework (Grade 9 to Grade 12)

- **Grade 9 (Explore & Experiment):** Try 3 different academic interests, start coding or practical building, and contribute to beginner open-source projects.
- **Grade 10 (Pick Your Core Spike):** Choose your singular focus area; compete in national olympiads (Mathematics, Physics, Informatics) or science fairs.

- **Grade 11 (Execute & Validate):** Build and launch a complete project; publish verified research in indexed preprints; file provisional patent applications.
- **Grade 12 (Polish & Apply):** Draft authentic college essays with DeepSeek V4 Flash; optimize Early Action (EA) applications and apply for merit scholarships.
- **Fake Journal Protection:** Automatically checks research journals against **Beall's List of Predatory Publishers** to prevent students from wasting money on scam "pay-to-publish" journals that damage their college credibility.

SECTION 12 · DATA PIPELINES & INGESTION CONNECTORS

Data Ingestion Infrastructure & External API Integration Contracts

To eliminate manual data entry, guarantee empirical freshness, and support the platform's zero-hardcoding mandate, IndiaLens operates an automated ingestion mesh comprising 18 specialized data connectors. Orchestrated via weekly and monthly Apache Airflow DAGs, these connectors continuously extract, normalize, validate, and persist economic, educational, and compensation telemetry.

💡 IN PLAIN ENGLISH · OUR AUTOMATED TRUTH-DETECTION MESH

College websites and brochures are filled with marketing exaggerations. To ensure students only see 100% verified facts, **IndiaLens runs 18 automated digital connectors** that continually pull official records directly from government portals, audited salary databases, and live admission counters. If a college claims their average salary mysteriously doubled overnight without proof, our quarantine gate automatically flags and blocks the fake number.

12.1 Detailed Specification of the 18 External Ingestion Connectors

API-01 · National Institutional Ranking Framework (NIRF)

Reliability: 0.95 (Govt Registry)

ENDPOINT & INGESTION PROTOCOL

`https://www.nirfindia.org/api/reports` — Annual batch PDF download via Crawl4AI + weekly portal scraper for mid-cycle circular amendments.

DOWNSTREAM CONSUMING ENGINE

Establishes baseline institutional tier classifications, anchors the primary historical median wage, and calculates faculty-to-student ratios for the composite scorecard.

EXTRACTED SCHEMA & FIELDS

`approved_intake`, `actual_student_strength`, `median_salary_graduating`, `graduating_batch_size`, `students_placed_count`, `phd_students_graduated`, `capital_expenditure_inr`, `operational_expenditure_inr`.

VALIDATION & QUARANTINE GATE

Flagged if reported median salary jumps by >35% year-over-year without corresponding audited NIRF financial disclosure notes.

API-02 · AmbitionBox Enterprise Compensation Registry

Reliability: 0.85 (Verified Users)

ENDPOINT & INGESTION PROTOCOL

`https://www.ambitionbox.com/api/v2/salaries` — Automated headless Playwright scrape cluster executing weekly on Sundays at 03:00 IST.

DOWNSTREAM CONSUMING ENGINE

Directly trains the Multi-Horizon XGBoost Quantile Regressors, populates the Work-Life Balance Degradation vector (V_8), and computes historical Salary Volatility (V_2).

EXTRACTED SCHEMA & FIELDS

`company_name`, `job_title`, `experience_band` (1Y, 3Y, 5Y, 7Y, 10Y+), `salary_min`, `salary_median`, `salary_max`, `work_life_balance_rating`, `job_security_rating`.

VALIDATION & QUARANTINE GATE

Requires a minimum sample size of $N \geq 15$ verified employee submissions per experience band before updating live percentile distributions.

API-03 · Naukri.com Enterprise Hiring Demand Feed

Reliability: 0.85 (Live Market Feed)

ENDPOINT & INGESTION PROTOCOL

`https://recruiter.naukri.com/api/v1/jobs/feed` — Weekly Playwright cluster ingesting 10,000+ active job postings across top-12 Indian metropolitan hubs.

DOWNSTREAM CONSUMING ENGINE

Computes the Credential Inflation index (V_4), Spatial Gini Concentration coefficient (V_5), and powers the Course Marketplace Skill-Gap matching algorithm.

EXTRACTED SCHEMA & FIELDS

`job_id`, `job_title`, `hiring_company`, `city`, `state`, `min_exp_years`, `max_exp_years`, `salary_min_inr`, `salary_max_inr`, `required_skills_array`, `job_function`, `posted_at`.

VALIDATION & QUARANTINE GATE

Deduplicates job postings using exact hash matching on `(company, title, city, salary_min)` within 30-day rolling windows.

API-04 · Indeed India Employer & Job Posting Registry

Reliability: 0.82 (Cross-Validation)

ENDPOINT & INGESTION PROTOCOL

`https://api.indeed.com/v2/jobsearch` — Weekly automated REST publisher API pulls.

DOWNSTREAM CONSUMING ENGINE

Provides independent cross-validation for Naukri job volumes, isolating localized hiring surges from broad macroeconomic market shifts.

EXTRACTED SCHEMA & FIELDS

`requisition_volume`, `posted_compensation_range`, `employer_response_latency`, `regional_hiring_density`, `qualification_requirements`.

VALIDATION & QUARANTINE GATE

Anomalies flagged if Naukri and Indeed job volume deltas exceed 40% divergence for identical Standard Occupational Classification (SOC) codes.

API-05 · MoSPI Periodic Labour Force Survey (PLFS)

Reliability: 0.98 (Sovereign Survey)

ENDPOINT & INGESTION PROTOCOL

`https://mospi.gov.in/data/plfs/quarterly` — Quarterly automated ingestion of raw national microdata CSV files via Python automated ETL.

DOWNSTREAM CONSUMING ENGINE

Calculates the non-degree counterfactual opportunity cost (OC_{PLFS}) across all Indian states and determines state-specific baseline wage floors.

EXTRACTED SCHEMA & FIELDS

`state_code`, `education_level_code`, `labor_force_participation_rate`, `worker_population_ratio`, `quarterly_earnings_inr`, `sectoral_unemployment_rate`.

VALIDATION & QUARANTINE GATE

Validated against historical state Gross State Domestic Product (GSDP) growth curves; outliers re-indexed using national Consumer Price Index (CPI).

API-06 · Reserve Bank of India (RBI) DBIE Macroeconomic API

Reliability: 0.99 (Central Bank)

ENDPOINT & INGESTION PROTOCOL

`https://dbie.rbi.org.in/DBIE/dbie.rbi?site=home` — Monthly DBIE automated SOAP/REST data pull.

DOWNSTREAM CONSUMING ENGINE

Dynamically calibrates the master real discount rate ($r(t)$), updates the Industry Cyclicity vector (V_3), and sets floating loan drift parameters in Monte Carlo models.

EXTRACTED SCHEMA & FIELDS

`policy_repo_rate`, `cpi_combined_inflation`, `wpi_inflation`, `gdp_growth_rate`, `mclr_median_lending_rate`, `rbi_klems_sector_gva_series`.

VALIDATION & QUARANTINE GATE

Immediate production cache invalidation triggered within 60 seconds of any Monetary Policy Committee (MPC) repo rate announcement.

API-07 · World Bank International Comparison Program (ICP)

Reliability: 0.95 (Multilateral)

ENDPOINT & INGESTION PROTOCOL

`https://api.worldbank.org/v2/country/IND/indicator/PA.NUS.PPP` — Quarterly REST JSON API synchronization.

DOWNSTREAM CONSUMING ENGINE

Normalizes cross-border compensation benchmarks on global explorer pages, displaying true localized purchasing power rather than deceptive nominal exchange rates.

EXTRACTED SCHEMA & FIELDS

`country_iso`, `ppp_conversion_factor_lcu_per_usd`, `price_level_index`, `gdp_ppp_current_intl_dollar`. Current benchmark: ~₹23.10 INR / USD PPP.

VALIDATION & QUARANTINE GATE

Cross-referenced with IMF World Economic Outlook (WEO) datasets prior to committing to production macro tables.

API-08 · Reddit Developer Sentiment & Qualitative NLP Pipeline Reliability: 0.70 (Crowdsourced)

ENDPOINT & INGESTION PROTOCOL

<https://oauth.reddit.com/r/developersIndia> — Weekly PRAW OAuth ingestion across subreddits: *r/developersIndia*, *r/india*, *r/CATpreparation*, *r/JEENETards*.

DOWNSTREAM CONSUMING ENGINE

Provides unstructured ground-level compensation signals and acts as an early-warning anomaly detector for placement scams at private institutions.

EXTRACTED SCHEMA & FIELDS

`post_id`, `subreddit`, `raw_text`, `extracted_salary_inr`, `extracted_company`, `extracted_yoe`, `sentiment_polarity`, `nlp_confidence_score`.

VALIDATION & QUARANTINE GATE

Fine-tuned RoBERTa Named Entity Recognition (NER) pipeline; posts with NLP confidence < 0.75 are discarded automatically.

API-09 · Internshala Early-Career Internship Ingestion Reliability: 0.80 (Early-Career)

ENDPOINT & INGESTION PROTOCOL

<https://internshala.com/api/v1/internships/feed> — Weekly headless crawler.

DOWNSTREAM CONSUMING ENGINE

Anchors the Year-1 entry-level stipend and pre-placement offer probability distributions for undergraduate tier-2/3 institutions.

EXTRACTED SCHEMA & FIELDS

`role_title`, `stipend_monthly_inr`, `duration_months`, `ppo_offered_flag`, `required_skills`, `applicant_volume`.

VALIDATION & QUARANTINE GATE

Excludes unpaid or performance-variable commission postings from the stipend statistical engine.

API-10 · PayScale India Compensation Registry Reliability: 0.82 (Verified)

ENDPOINT & INGESTION PROTOCOL

<https://api.payscale.com/v2/compensation/india> — Monthly automated scrape.

DOWNSTREAM CONSUMING ENGINE

Validates 10-year and 15-year salary horizons in the BiLSTM deep trajectory network to prevent long-term extrapolation drift.

EXTRACTED SCHEMA & FIELDS

`degree_type`, `major_discipline`, `experience_years`, `p10_comp`, `p50_comp`, `p90_comp`, `annual_bonus_inr`.

VALIDATION & QUARANTINE GATE

Re-indexes historical profiles using actual CPI inflation if reported compensation lags by >24 months.

API-11 · US Bureau of Labor Statistics (BLS) OEWS Database Reliability: 0.98 (US Govt)

ENDPOINT & INGESTION PROTOCOL

<https://api.bls.gov/publicAPI/v2/timeseries/data/> — Annual Occupational Employment & Wage Statistics JSON API ingestion.

DOWNSTREAM CONSUMING ENGINE

Establishes the empirical baseline gross salary curves for all US master's and undergraduate degree valuation models.

EXTRACTED SCHEMA & FIELDS

`soc_code`, `occupation_title`, `annual_mean_wage_usd`, `p10_wage_usd`, `p25_wage_usd`, `p75_wage_usd`, `p90_wage_usd`, `10_year_job_outlook_pct`.

VALIDATION & QUARANTINE GATE

Adjusted annually following the official May release of BLS national wage tables.

API-12 · Levels.fyi Global High-Tech Compensation Feed Reliability: 0.92 (Verified Paystubs)

ENDPOINT & INGESTION PROTOCOL

<https://api.levels.fyi/v1/salaries> — Weekly JSON sync.

DOWNSTREAM CONSUMING ENGINE

Models the extreme upper-tail (P_{90}) wage upside for Tier-1 domestic institutions (IIT/BITS) and elite global computer science programs (CMU/Stanford/MIT).

EXTRACTED SCHEMA & FIELDS

`company_name`, `level_title` (e.g., L3/L4/L5, SDE-1/2/3), `base_salary_usd`, `stock_grant_annual_usd`, `annual_bonus_usd`, `total_compensation_usd`, `location_metro`.

VALIDATION & QUARANTINE GATE

Requires verified employee offer letter or W-2/Form-16 submission flags before inclusion in the compensation sample.

API-13 · Glassdoor Global Compensation & Satisfaction Feed

Reliability: 0.80 (Aggregated)

ENDPOINT & INGESTION PROTOCOL

`https://api.glassdoor.com/v1/reviews` — Monthly automated connector.

DOWNSTREAM CONSUMING ENGINE

Cross-validates city-level rent and living cost assumptions for foreign master's programs in Europe, the UK, and North America.

EXTRACTED SCHEMA & FIELDS

`metro_area`, `cost_of_living_score`, `median_rent_monthly_usd`, `culture_rating`, `career_opportunities_rating`, `interview_difficulty_score`.

VALIDATION & QUARANTINE GATE

Sanity checked against Numbeo Cost of Living Indexes to eliminate localized survey biases.

API-14 · DeepSeek V4 Flash Cognitive Memory & Fiduciary Engine

Reliability: 0.98 (Frontier LLM)

ENDPOINT & INGESTION PROTOCOL

`https://api.deepseek.com/v1/chat/completions` (Model: `deepseek-v4-flash`) — Real-time streaming WebSocket/REST API per user turn.

DOWNSTREAM CONSUMING ENGINE

Drives the interactive Student Intelligence Copilot, executes Pass 1 constraint formulation, authors Socratic portfolio essays, and synthesizes Pass 2 downside-first fiduciary memorandums.

EXTRACTED SCHEMA & FIELDS

`streaming_tokens`, `prompt_cache_hit_tokens`, `tool_calls` (for Path DAG mutation), `finish_reason`, `usage`.

VALIDATION & QUARANTINE GATE

Strictly enforces zero numerical hallucination; answers failing verification against the `actuarial_ml_outputs` packet are regenerated before reaching the UI.

API-15 · Perplexity Sonar Live Research & Grounding Gateway

Reliability: 0.96 (Live Search Index)

ENDPOINT & INGESTION PROTOCOL

`https://api.perplexity.ai/chat/completions` (Model: `sonar`) — Real-time async REST query with native search indexing.

DOWNSTREAM CONSUMING ENGINE

Powers the Fork-and-Join pipeline; extracts real-time tuition revisions, cutoffs, and regulatory alerts; splits numerical params to Python ML and qualitative context to DeepSeek.

EXTRACTED SCHEMA & FIELDS

`content`, `citations` (array of verified URLs), `search_results` (structured title, URL, snippet), `query_plan`.

VALIDATION & QUARANTINE GATE

Strict domain indexing prioritizes official education and government portals (`.gov.in`, `.nic.in`, `.ac.in`, `.edu`); cached in Redis with a 72-hour TTL.

API-16 · Jina Reader Markdown Extraction Pipeline

Reliability: 0.92 (Document Parser)

ENDPOINT & INGESTION PROTOCOL

`https://r.jina.ai/{target_url}` — On-demand async HTTP GET request.

DOWNSTREAM CONSUMING ENGINE

Converts messy, unstructured university circulars, fee schedules, and curriculum PDFs into clean markdown for automated LLM entity extraction.

EXTRACTED SCHEMA & FIELDS

`clean_markdown_text`, `parsed_tables_json`, `page_title`, `http_status_code`.

VALIDATION & QUARANTINE GATE

Ensures non-empty markdown output (>200 characters); falls back to Crawl4AI on extraction failures.

API-17 · Crawl4AI Distributed Crawler Engine

Reliability: 0.94 (Internal Cluster)

ENDPOINT & INGESTION PROTOCOL

`http://localhost:11235/crawl` — Batch async web crawler cluster running on local Docker infrastructure.

DOWNSTREAM CONSUMING ENGINE

Primary crawler for JavaScript-heavy state counseling portals (JoSAA, CSAB, JAC Delhi, WBJEE, KEAM, COMEDK).

EXTRACTED SCHEMA & FIELDS

`raw_html`, `rendered_dom_tree`, `javascript_evaluated_tables`, `downloaded_pdf_binaries`.

VALIDATION & QUARANTINE GATE

Automatic retry with rotating residential proxies upon encountering rate limits or Cloudflare challenges.

ENDPOINT & INGESTION PROTOCOL

Coursera Partner API / Impact.com Gateway / Direct OAuth — Real-time server-to-server webhook callbacks.

DOWNSTREAM CONSUMING ENGINE

Synchronizes the Course Marketplace catalog, validates coupon codes, tracks attribution cookies, and reconciles commercial revenues.

EXTRACTED SCHEMA & FIELDS

course_catalog_json, click_uuid, conversion_time_stamp, commission_payout_usd, student_retention_f_lag.

VALIDATION & QUARANTINE GATE

SHA-256 HMAC cryptographic signature validation on all incoming webhook payloads.

SECTION 13 · RELATIONAL PERSISTENCE LAYER

Production Relational Database Schema (PostgreSQL 15+)

The persistence layer is implemented in PostgreSQL 15+ utilizing native JSONB document storage for dynamic graph structures, strict foreign-key relational integrity, and GiST/GIN indexes for sub-millisecond query latency.

IN PLAIN ENGLISH · THE MASTER DIGITAL HEALTH CHART FOR YOUR CAREER

Why 21 tables? Think of this database like a high-precision medical health record, but built for your educational future. It securely stores your personal quiz strengths, connects real-world college fees, tracks 20-year alumni salary curves, stress-tests bank loan terms, and monitors AI job threats—allowing IndiaLens to deliver personalized, lightning-fast career answers in milliseconds.

13.1 Complete Specifications for All 21 Core Database Tables

Table 01 · colleges

Primary Relational Directory

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, short_name VARCHAR(64) NOT NULL, full_name VARCHAR(255) NOT NULL, state VARCHAR(64) NOT NULL, city VARCHAR(64) NOT NULL, tier SMALLINT CHECK (tier IN (1,2,3)), college_type VARCHAR(32) (IIT/NIT/BITS/State/Private), naac_grade VARCHAR(8), nirf_rank INT, established_year INT, website_url TEXT, is_active BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Master index of all recognized domestic tertiary institutions. Indexed on state, tier, nirf_rank.

Table 02 · degrees

Standardized Degree Taxonomy

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, short_name VARCHAR(32) NOT NULL, full_name VARCHAR(128) NOT NULL, field VARCHAR(64) NOT NULL, level VARCHAR(16) CHECK (level IN ('UG','PG','PhD','Diploma')), duration_years INT NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

Taxonomy of standardized academic credentials (e.g., B.Tech Computer Science, MBBS, MBA Finance).

Table 03 • programs

Core Unit of Econometric Evaluation

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, college_id UUID REFERENCES colleges(id), degree_id UUID REFERENCES degrees(id), annual_tuition_inr NUMERIC(12,2) NOT NULL, total_seats INT, cutoff_rank_low INT, cutoff_rank_high INT, is_active BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Represents a specific degree offering at a specific college. Unique constraint on (college_id, degree_id).

Table 04 • data_points

Immutable Scraping Audit Ledger

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, program_id UUID REFERENCES programs(id), scrape_run_id UUID REFERENCES scrape_runs(id), field_name VARCHAR(64) NOT NULL, raw_value TEXT, parsed_value NUMERIC(14,4), unit VARCHAR(32), source_url TEXT, scraped_at TIMESTAMPTZ NOT NULL, is_current BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Complete historical audit trail for every scraped datapoint. Strict partial unique index on (program_id, field_name) WHERE is_current=TRUE.

Table 05 • roi_scores

Computed Multi-Dimensional Index

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, program_id UUID REFERENCES programs(id), model_version VARCHAR(32) REFERENCES model_versions(version_tag), composite_score NUMERIC(5,2) NOT NULL, financial_roi_pct NUMERIC(8,2), risk_score NUMERIC(5,2), optionality_score NUMERIC(5,2), mobility_score NUMERIC(5,2), satisfaction_score NUMERIC(5,2), network_score NUMERIC(5,2), ci_low NUMERIC(5,2), ci_high NUMERIC(5,2), is_current BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Stores the 0-100 executive composite ROI score and its component dimensions for immediate dashboard query rendering.

Table 06 • salary_trajectories

Multi-Horizon Wage Percentiles

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, program_id UUID REFERENCES programs(id), model_version VARCHAR(32) REFERENCES model_versions(version_tag), year_number INT NOT NULL (1, 2, 3, 5, 7, 10, 15, 20), p10_inr NUMERIC(12,2), p25_inr NUMERIC(12,2), p50_inr NUMERIC(12,2), p75_inr NUMERIC(12,2), p90_inr NUMERIC(12,2), ppp_usd_p50 NUMERIC(12,2), sample_size INT, is_current BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Powers the Recharts wage dispersion fan charts on report dashboards across multiple career horizons.

Table 07 • risk_indicators

8-Vector AI Risk Surface

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, program_id UUID REFERENCES programs(id), model_version VARCHAR(32) REFERENCES model_versions(version_tag), ai_automation_prob NUMERIC(5,4), salary_volatility NUMERIC(5,4), industry_cyclicality NUMERIC(5,4), credential_inflation NUMERIC(5,4), geographic_concentration NUMERIC(5,4), regulatory_risk NUMERIC(5,4), physical_hazard NUMERIC(5,4), wlb_degradation NUMERIC(5,4), jss_score NUMERIC(5,2), ai_risk_label VARCHAR(64), is_current BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Stores the 8-vector risk surface scores and computed Job Security Score (JSS) per degree program.

Table 08 • placement_data

Verified Placement Audit

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, program_id UUID REFERENCES programs(id), academic_year VARCHAR(16) NOT NULL, placement_rate_pct NUMERIC(5,2), median_salary_inr NUMERIC(12,2), average_salary_inr NUMERIC(12,2), highest_salary_inr NUMERIC(12,2), companies_visited INT, ppo_count INT, source VARCHAR(64), is_current BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Historical placement audit records ingested directly from NIRF, AISHE, and official Right to Information (RTI) filings.

Table 09 • cost_data

Granular Capital Expenditure

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, program_id UUID REFERENCES programs(id), total_tuition_inr NUMERIC(12,2) NOT NULL, hostel_living_inr NUMERIC(12,2) NOT NULL, exam_prep_costs_inr NUMERIC(12,2) DEFAULT 0, opportunity_cost_inr NUMERIC(12,2) NOT NULL, total_cost_of_degree NUMERIC(12,2) NOT NULL, source_year INT NOT NULL, is_current BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Complete financial capital expenditure breakdown used in master NPV and Monte Carlo loan simulations.

Table 10 • student_reports

Persisted Analytical Dossiers

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, token VARCHAR(64) UNIQUE NOT NULL, profile_data JSONB NOT NULL, results_data JSONB NOT NULL, model_version VARCHAR(32) NOT NULL, generated_at TIMESTAMPTZ NOT NULL, viewed_count INT DEFAULT 0, expires_at TIMESTAMPTZ NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

Stores generated actuarial reports accessible via secure tokenized URLs (/report/:token) with 90-day persistence.

Table 11 • personal_intelligence

Longitudinal Student Memory & Path DAG

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, token VARCHAR(64) UNIQUE NOT NULL, profile_data JSONB NOT NULL, intelligence JSONB NOT NULL, path_graph JSONB NOT NULL, citations JSONB DEFAULT '[]'::jsonb, grounded BOOL DEFAULT TRUE, created_at TIMESTAMPTZ NOT NULL, updated_at TIMESTAMPTZ NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

Stateful memory store for the DeepSeek V4 Flash Copilot, holding the dynamic Path DAG, conversation turns, Sonar citations, and rejected option logs.

Table 12 • scrape_runs

Airflow Orchestration Run Logs

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, source_name VARCHAR(64) NOT NULL, started_at TIMESTAMPTZ NOT NULL, completed_at TIMESTAMPTZ, status VARCHAR(16) NOT NULL CHECK (status IN ('running','success','partial','failed')), records_scraped INT DEFAULT 0, records_updated INT DEFAULT 0, records_flagged INT DEFAULT 0, error_log TEXT.

OPERATIONAL PURPOSE & CONSTRAINTS

Operational health telemetry tracking scraping pipelines and triggering automated retraining upon successful execution.

Table 13 • anomalies

Data Quarantine Gate

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, program_id UUID REFERENCES programs(id), scrape_run_id UUID REFERENCES scrape_runs(id), field_name VARCHAR(64) NOT NULL, prior_value NUMERIC(14,4), new_value NUMERIC(14,4), delta_pct NUMERIC(6,2) NOT NULL, status VARCHAR(16) DEFAULT 'pending' CHECK (status IN ('pending','accepted','rejected','auto_accepted')), reviewed_by VARCHAR(64), reviewed_at TIMESTAMPTZ.

OPERATIONAL PURPOSE & CONSTRAINTS

Quarantines any scraped metrics exhibiting >25% variation until manually audited or statistically validated.

Table 14 • model_versions

ML Model Champion / Challenger Registry

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, version_tag VARCHAR(32) UNIQUE NOT NULL, is_live BOOL DEFAULT FALSE, trained_at TIMESTAMPTZ NOT NULL, training_records INT NOT NULL, mse NUMERIC(10,4) NOT NULL, r2_score NUMERIC(6,4) NOT NULL, mape_salary NUMERIC(6,4) NOT NULL, recall_at_5 NUMERIC(6,4), changelog TEXT, artifact_path TEXT NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

Registry governing XGBoost and BiLSTM model weights; enforces automated promotion gates before updating live inference engines.

Table 15 • job_postings

Real-Time Hiring Demand Stream

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, scrape_run_id UUID REFERENCES scrape_runs(id), title VARCHAR(255) NOT NULL, company VARCHAR(255) NOT NULL, city VARCHAR(64) NOT NULL, state VARCHAR(64), experience_min INT, experience_max INT, salary_min_inr NUMERIC(12,2), salary_max_inr NUMERIC(12,2), skills_required TEXT[] NOT NULL, job_function VARCHAR(64), posted_at TIMESTAMPTZ NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

Ingested market hiring demand feed powering skill-gap vector calculations and spatial concentration metrics.

Table 16 • reddit_extractions

Crowdsourced Compensation NLP

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, scrape_run_id UUID REFERENCES scrape_runs(id), post_id VARCHAR(32) UNIQUE NOT NULL, subreddit VARCHAR(64) NOT NULL, title TEXT NOT NULL, body_snippet TEXT, extracted_salary NUMERIC(12,2), extracted_company VARCHAR(128), extracted_role VARCHAR(128), extracted_yoe NUMERIC(4,1), nlp_confidence NUMERIC(4,3) NOT NULL, post_date TIMESTAMPTZ NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

Parsed community salary disclosures extracted via RoBERTa NER models for sentiment and fraud anomaly flagging.

Table 17 • macro_indicators

Sovereign Macroeconomic State

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, scrape_run_id UUID REFERENCES scrape_runs(id), indicator_key VARCHAR(64) UNIQUE NOT NULL, parsed_value NUMERIC(14,4) NOT NULL, unit VARCHAR(32) NOT NULL, source_url TEXT NOT NULL, updated_at TIMESTAMPTZ NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

Single source of truth for global economic constants: RBI Repo Rate, CPI, PPP factor (~₹23.10), and USD/EUR FX drift parameters.

Table 18 • course_marketplace

Curated Upskilling Catalog

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, course_title VARCHAR(255) NOT NULL, provider VARCHAR(64) NOT NULL, category VARCHAR(64) NOT NULL, affiliate_url TEXT NOT NULL, price_inr NUMERIC(10,2) NOT NULL, duration_hours INT NOT NULL, skill_tags TEXT[] NOT NULL, career_paths TEXT[] NOT NULL, ai_resilience_score NUMERIC(4,3) NOT NULL, commission_rate_pct NUMERIC(5,2) NOT NULL, is_active BOOL DEFAULT TRUE.

OPERATIONAL PURPOSE & CONSTRAINTS

Indexed inventory of accredited upskilling courses available for dynamic matching via the Course Match Score formula.

Table 19 • course_impressions

Marketplace Monetization Telemetry

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, student_token VARCHAR(64) NOT NULL, course_id UUID REFERENCES course_marketplace(id), match_score NUMERIC(4,3) NOT NULL, shown_at TIMESTAMPTZ NOT NULL, clicked BOOL DEFAULT FALSE, converted BOOL DEFAULT FALSE, revenue_inr NUMERIC(10,2) DEFAULT 0.

OPERATIONAL PURPOSE & CONSTRAINTS

Monitors recommendation Click-Through Rates (CTR), conversion yields, and affiliate revenue attribution across student cohorts.

Table 20 • global_programs

International Institutional Offerings

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, university_name VARCHAR(255) NOT NULL, country VARCHAR(64) NOT NULL, city VARCHAR(64) NOT NULL, degree_name VARCHAR(128) NOT NULL, major VARCHAR(128) NOT NULL, is_stem_designated BOOL NOT NULL, annual_tuition_usd NUMERIC(10,2) NOT NULL, living_cost_annual_usd NUMERIC(10,2) NOT NULL, scholarship_probability NUMERIC(4,3), assistantship_probability NUMERIC(4,3), median_salary_usd_y1 NUMERIC(10,2) NOT NULL, visa_type VARCHAR(32) NOT NULL, visa_survival_prob NUMERIC(4,3) NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

International degree database evaluated via cross-border Net Present Value and H-1B/Blue Card survival mathematics.

Table 21 • portfolio_profiles

Global Admissions Spike State

COLUMNS & DATA TYPES

id UUID PRIMARY KEY, student_token VARCHAR(64) UNIQUE NOT NULL, spike_domain VARCHAR(128) NOT NULL, target_universities JSONB NOT NULL, activities JSONB NOT NULL, research_artifacts JSONB NOT NULL, essay_drafts JSONB NOT NULL, milestone_schedule JSONB NOT NULL, spike_authenticity_score NUMERIC(5,2) NOT NULL, updated_at TIMESTAMPTZ NOT NULL.

OPERATIONAL PURPOSE & CONSTRAINTS

Maintains the student's multi-year holistic admissions portfolio, extracurricular artifacts, research preprints, and Common App essay drafts.

SECTION 14 • USER EXPERIENCE & INTERFACE ARCHITECTURE

Page-by-Page UI/UX Architecture & Screen Specifications

The visual and operational design language of IndiaLens rejects the colorful, gamified aesthetic common in legacy EdTech. Instead, it adopts a high-density, professional **Financial Terminal aesthetic** (combining the data density of Bloomberg with modern, accessible web ergonomics) to instill immediate institutional credibility and trust.

💡 IN PLAIN ENGLISH · HOW STUDENTS & PARENTS EXPERIENCE INDIALENS

We deliberately rejected childish, gamified cartoons and distracting college advertisements. Instead, **IndiaLens is built like a clean, high-precision financial dashboard** that anyone can use in seconds. You get clear wage charts showing your likely salary growth, simple sliders to test what happens if bank interest rates rise, and straightforward color-coded cards that highlight college risks without confusion.

14.1 Comprehensive Screen Specifications Across All 12 Routes

Route 01 · / (Master Sovereign Landing)

Entryway & Sovereign Proposition

KEY VISUAL COMPONENTS & WIDGETS

Live Indexed Programs Ticker (3,400+ programs across India/Global), Interactive 3-College ROI Card Carousel, NIRF vs. IndiaLens Asymmetric Contrast Table, Real-time Macroeconomic Ticker Bar (RBI Repo Rate, MoSPI Baseline Earnings, US STEM H-1B Odds).

BACKEND ENDPOINTS & API CALLS

`GET /api/v2/analytics/summary`, `GET /api/v2/colleges/featured`, `GET /api/v2/macro/live-ticker`.

USER INTERACTIONS & WORKFLOWS

Hero CTA initiates the 4-Step Intake Wizard; an interactive fee-vs-salary slider allows visitors to test programmatic career arbitrage instantly without logging in.

STATE MANAGEMENT & PERFORMANCE TARGET

Static Site Generation (SSG) with 60-second Incremental Static Regeneration (ISR); sub-1.2s First Contentful Paint (FCP) on mobile 4G.

Route 02 · /analyze (Multi-Stage Intake Wizard)

Psychometric & Constraints Diagnostic

KEY VISUAL COMPONENTS & WIDGETS

4-Stage Progressive Stepper: Step 1 (Academic Baseline: Board marks, JEE/CAT/SAT mock scores, category, home state), Step 2 (Financial Constraints: Maximum family budget, loan appetite, monthly EMI ceiling), Step 3 (6-8 Question 3PL IRT CAT Micro-Dilemma Diagnostic), Step 4 (Stochastic Monte Carlo Simulation Loader).

BACKEND ENDPOINTS & API CALLS

`POST /api/v2/psychometrics/next-item`, `POST /api/v2/analyze/generate-dossier`.

USER INTERACTIONS & WORKFLOWS

Adaptive CAT questions load asynchronously based on previous Bayesian trait convergence; real-time Fisher Information progress indicator; generates a secure 64-character UUID token and redirects to the intelligence dossier.

STATE MANAGEMENT & PERFORMANCE TARGET

Zustand client-side state store persisting partial responses in `localStorage`; zero data loss on accidental page reload.

Route 03 •
/report/:token
(The Actuarial
Intelligence
Dossier)

Primary Actuarial Artifact

KEY VISUAL COMPONENTS & WIDGETS

Executive Composite ROI Score Ring (0–100), 20-Year Wage Dispersion Fan Chart (Recharts rendering $P_{10}, P_{25}, P_{50}, P_{75}, P_{90}$ curves), Monte Carlo Net Worth Curves, 8-Vector AI Risk Radar Chart, Loan Solvency Stress Barometer (Default Probability %), Grounded DeepSeek V4 Flash Fiduciary Synthesis Memorandum (Paired with Sonar Grounding).

BACKEND ENDPOINTS & API CALLS

`GET /api/v2/reports/{token}`, `POST /api/v2/reports/{token}/re-simulate`, `POST /api/v2/reports/{token}/export-pdf`.

USER INTERACTIONS & WORKFLOWS

Users toggle between P10 (conservative) and P90 (optimistic) career scenarios; dynamic sliders recalculate floating loan EMIs in real-time; single-click action launches the Copilot to interrogate specific report assumptions.

STATE MANAGEMENT & PERFORMANCE TARGET

Server-Side Rendered (SSR) with React Query caching; client-side re-simulation updates in < 150ms.

Route 04 • /explore (Institutional Asset Screener)

Quantitative Search & Screening

KEY VISUAL COMPONENTS & WIDGETS

Multi-Faceted Filter Sidebar (Tier 1/2/3, Academic Field, State Domicile Quota, Maximum Annual Tuition, Minimum Composite Score, Maximum AI Disruption Risk), High-Density Data Grid with sortable columns, Embedded Mini Sparkline Wage Trajectories.

BACKEND ENDPOINTS & API CALLS

`GET /api/v2/programs/search`, `GET /api/v2/filters/facets`.

USER INTERACTIONS & WORKFLOWS

Instant 150ms debounced client-side filtering; multi-select checkbox allows users to stage up to 3 programs for comparative evaluation.

STATE MANAGEMENT & PERFORMANCE TARGET

URL query parameter synchronization; TanStack Table virtualization supporting 5,000+ rows at 60fps.

Route 05 • /college/:id (Programmatic Deep-Dive Terminal)

Institutional Audit Terminal

KEY VISUAL COMPONENTS & WIDGETS

Institutional Scorecard, Verified RTI Placement Audit Table (Median vs. Highest vs. Batch Placement Rate), Historical Closing Rank Cutoff Trends, Granular 8-Vector AI Risk Breakdown with Skill Moat Recommendations, Embedded Grounded Advisor Widget (DeepSeek + Sonar).

BACKEND ENDPOINTS & API CALLS

`GET /api/v2/colleges/{id}`, `GET /api/v2/programs/{id}/risk-surface`, `GET /api/v2/programs/{id}/cutoffs`.

USER INTERACTIONS & WORKFLOWS

Deep branch-level placement inspection; interactive loan amortization scheduler; direct action button to add program as a target node in the student's personal Path DAG.

STATE MANAGEMENT & PERFORMANCE TARGET

Fully static pre-rendering for top-500 institutions with dynamic client-side hydration for user-conditioned fit overlays.

Route 06 • /compare (Side-by-Side Econometric Comparator)

Head-to-Head Valuation

KEY VISUAL COMPONENTS & WIDGETS

Tri-Column Head-to-Head Comparative Matrix: Net Present Value Differential (ΔNPV), Breakeven Horizon Gap (ΔT_{BEP}), Downside Risk Overlay (P10 vs P10), Student Psychometric Fit Delta (ΔF_{it}).

BACKEND ENDPOINTS & API CALLS

`POST /api/v2/programs/compare`, `POST /api/v2/reports/export-comparison`.

USER INTERACTIONS & WORKFLOWS

Dynamic visual highlighting of mathematical winners across every econometric dimension; exports formal comparison summary as a branded PDF memorandum.

STATE MANAGEMENT & PERFORMANCE TARGET

Zustand comparative tray synced to browser sessionStorage.

Route 07 · /advisor (Personalized Intelligence Copilot)

Grounded Generative Fiduciary

KEY VISUAL COMPONENTS & WIDGETS

Full-Screen Conversational Terminal, Active Context Drawer (displaying student constraints, IRT trait vector, and loan tolerance), Inline Grounded Citation Cards with verified source URLs, Socratic Counterargument Prompts.

BACKEND ENDPOINTS & API CALLS

WS /api/v2/ai/chat-stream, POST /api/v2/ai/mutate-dag, GET /api/v2/ai/context-stack.

USER INTERACTIONS & WORKFLOWS

Live streaming multi-turn conversation with DeepSeek V4 Flash (Cognitive Memory Core) & Perplexity Sonar (Epistemic Grounding Core); users can challenge advice, request alternative scenario simulations, and approve Path DAG updates directly via conversational actions.

STATE MANAGEMENT & PERFORMANCE TARGET

WebSocket streaming with optimistic UI updates; initial token latency < 400ms.

Route 08 · /methodology (Open Actuarial Ledger)

Transparency & Formula Auditing

KEY VISUAL COMPONENTS & WIDGETS

Full Mathematical Formula Documentation (typeset via KaTeX), Live Data Freshness Badges (displaying exact scrape timestamps per source connector), Model Weight Calibration Transparency Ledger.

BACKEND ENDPOINTS & API CALLS

GET /api/v2/meta/data-freshness, GET /api/v2/meta/model-versions.

USER INTERACTIONS & WORKFLOWS

Interactive formula inspector; allows economists and parents to verify mathematical variables and download replication datasets.

STATE MANAGEMENT & PERFORMANCE TARGET

Static documentation page with 100% SEO indexing and schema.org structured academic metadata.

Route 09 · /global (Global Degree Valuation Explorer)

Cross-Border Valuation Hub

KEY VISUAL COMPONENTS & WIDGETS

Country Parity Selector (US, UK, Germany, Canada, Singapore), STEM OPT Designation Filter, Currency & Purchasing Power Parity (PPP) Toggle, Spatial COLI Tax Net-Savings Comparison Matrix.

BACKEND ENDPOINTS & API CALLS

GET /api/v2/global/programs, POST /api/v2/global/cross-border-npv.

USER INTERACTIONS & WORKFLOWS

Cross-border valuation modeling; users can simulate foreign master's programs against domestic Tier-1 alternatives on an equalized net savings basis.

STATE MANAGEMENT & PERFORMANCE TARGET

Interactive client-side FX recalculations executed in < 50ms.

Route 10 · /portfolio-builder (Admissions Spike & Research Studio)

Holistic Portfolio Optimization

KEY VISUAL COMPONENTS & WIDGETS

Spike Authenticity Meter (A_{spike}), Action-Impact X-Y-Z Transformer, 4-Year Milestone Gantt Scheduler, Research Topic Clustering Engine, Predatory Journal Scanner (Beall's List Crosscheck).

BACKEND ENDPOINTS & API CALLS

POST /api/v2/portfolio/transform-xyz, POST /api/v2/portfolio/check-preprint, GET /api/v2/portfolio/profile.

USER INTERACTIONS & WORKFLOWS

Students draft activity bullets; DeepSeek V4 Flash provides real-time restructuring into high-impact X-Y-Z format; validates preprints against Beall's List.

STATE MANAGEMENT & PERFORMANCE TARGET

Rich text editor with autosave to PostgreSQL JSONB every 3 seconds.

Route 11 · /marketplace (Actuarial Upskilling Marketplace)

Curated Skill Acquisition

KEY VISUAL COMPONENTS & WIDGETS

Curated Course Grid, Match Score Badges ($M \geq 0.75$), Projected Salary Uplift Indicators ($\Delta \hat{S}_{Y1}$), Provider Badges (Coursera, edX, Unacademy), Category Filter Tabs.

BACKEND ENDPOINTS & API CALLS

`GET /api/v2/marketplace/recommended`, `POST /api/v2/marketplace/track-click`.

USER INTERACTIONS & WORKFLOWS

Clicking a course securely logs impression, signs affiliate tracking token, and redirects to partner provider with 30-day attribution.

STATE MANAGEMENT & PERFORMANCE TARGET

Client-side personalization mesh rendering matched cards in < 200ms.

Route 12 · /admin (Operational Telemetry Control Center)

Operations & Governance

KEY VISUAL COMPONENTS & WIDGETS

Airflow Ingestion Run Monitors, Scrape Quarantine Anomaly Queue, Model Registry (Champion vs. Challenger metrics), Revenue & Affiliate Conversion Telemetry.

BACKEND ENDPOINTS & API CALLS

`GET /api/v2/admin/telemetry`, `POST /api/v2/admin/promote-model`, `POST /api/v2/admin/resolve-anomaly`.

USER INTERACTIONS & WORKFLOWS

Internal operations dashboard for managing data anomalies, manually approving scraped fee updates, promoting ML models, and auditing affiliate payouts.

STATE MANAGEMENT & PERFORMANCE TARGET

Role-Based Access Control (RBAC) with MFA; internal audit logging on all database mutations.

SECTION 15 · COMMERCIALIZATION & FINANCIAL MODEL

Monetization Architecture, Unit Economics & Financial Projections

IndiaLens strictly rejects the predatory lead-generation advertising model. By maintaining complete fiduciary independence from university marketing budgets, IndiaLens monetizes through consumer premium analytics, high-intent curated upskilling affiliate commissions, and enterprise B2B actuarial data licensing.

💡 IN PLAIN ENGLISH · HONEST, TRANSPARENT BUSINESS ECONOMICS

Traditional college portals make money by selling your phone number to private colleges for ₹500 to ₹2,500. **IndiaLens operates on complete independence:** We never take advertising money from colleges. Students pay a modest fee (₹499 one-time or ₹199/month) for full access to 20-year salary graphs and loan stress simulations. In addition, education lenders and banks pay us to evaluate loan repayment safety. This ensures our only loyalty is to the student and family.

15.1 Multi-Tiered Feature & Paywall Matrix

TABLE 15.1 – COMMERCIAL TIERING & PAYWALL ARCHITECTURE

Platform Capability	Free Sovereign Tier (₹0)	Premium Student Pass (₹499 / ₹199 mo)	Institutional B2B License
Institutional Discovery & Search	Unlimited access to /explore and basic college profiles	Unlimited access with advanced econo-metric filtering	Custom white-labeled portal + bulk student roster export
Econometric Wage Modeling	Point Median Salary (Y1 and Y5) only	Full 20-Year Wage Dispersion Fan Charts ($P_{10}, P_{25}, P_{50}, P_{75}, P_{90}$)	Full cohort simulation API with cus-tom wage benchmarks
Debt Stress & Solvency Simulation	Basic static EMI calculation	Full 10,000-Path Monte Carlo Loan Default & Solvency Simulation	Institutional credit-underwriting risk scoring API
AI Risk & Disruption Surface	High-level qualitative risk label	Full 8-Vector Risk Radar, JSS Score, and 5Y/10Y Tail Layoff Probabilities	Enterprise workforce automation ex-posure dataset access
Personalized Copilot (DeepSeek + Sonar)	15 exploratory messages per ses-sion (sub-cent cached inference)	Unlimited multi-turn grounded fiduciary ad-visory sessions	Dedicated enterprise counselor workspace with audit logs
Global Degree ROI & Visa Math	Top-3 foreign university overview	Comprehensive cross-border NPV, STEM OPT math & COLI tax drag models	Full international mobility and schol-arship data feeds
Portfolio Builder & Spike Studio	1 activity X-Y-Z optimization preview	Full 4-Year Milestone Tracker, Research Incubator, and Common App Studio	School-wide extracurricular gover-nance and verification engine
Report Artifact Generation	Web view only (tokenized link)	Downloadable 16-page branded Actuarial Dossier (PDF)	Bulk branded institutional PDF report generation

15.2 Unit Economics & 3-Year Projections

TABLE 15.2 – 3-YEAR OPERATIONAL & FINANCIAL FORECAST

Operational Metric	Year 1 (Foundation)	Year 2 (Scaling)	Year 3 (Dominance)
Annual Active Free Users	450,000	2,200,000	7,500,000
Paid Premium Subscriptions (Conversion)	22,500 (5.0%)	154,000 (7.0%)	675,000 (9.0%)
Gross Premium Subscription Revenue	₹1.12 Cr	₹7.68 Cr	₹33.68 Cr
Course Marketplace Conversions	18,000	110,000	420,000
Gross Affiliate Revenue (₹1,200 avg comm)	₹2.16 Cr	₹13.20 Cr	₹50.40 Cr
Enterprise B2B Contracts (Coaching/NBFC)	15 (₹75 L)	65 (₹3.90 Cr)	220 (₹15.40 Cr)
Total Annual Gross Revenue	₹4.03 Cr	₹24.78 Cr	₹99.48 Cr
Customer Acquisition Cost (Blended CAC)	₹42 / user	₹28 / user	₹18 / user
Lifetime Value to CAC Ratio (LTV / CAC)	19.8x (Re-Indexed)	24.2x	34.5x

STRATEGIC ARCHITECTURAL MOAT • 90% AI OPEX REDUCTION

By replacing the monolithic Gemini 3.7 Flash stack (\$35/1,000 queries) with the **Fork-and-Join DeepSeek V4 Flash + Perplexity Sonar** architecture (\$3.20/1,000 queries blended), IndiaLens slashes direct inference OPEX by ~88%. This structural cost advantage allows us to triple the free-tier conversational allowance (from 5 to 15 turns per session) without burning runway, accelerating top-of-funnel viral acquisition while expanding Year 1 net operating margins to 84%.